

# Introduction to Scientific Programming in C++/Fortran2003

Victor Eijkhout

2017



# Contents

## I Introduction 9

1	<b>Introduction</b>	11
1.1	<i>Programming and computational thinking</i>	11
1.1.1	History	11
1.1.2	Computational thinking	13
1.1.3	Hardware	14
1.1.4	Algorithms	14
1.2	<i>About the choice of language</i>	14
1.3	<i>Further reading</i>	15
2	<b>Warming up</b>	17
2.1	<i>Programming environment</i>	17
2.1.1	Language support in your editor	17
2.2	<i>Compiling</i>	18
2.2.1	Compiling C++	18
2.2.2	Compiling Fortran	19
3	<b>Teachers guide</b>	21
3.1	<i>Justification</i>	21
3.2	<i>Timeline for a C++/F03 course</i>	21
3.2.1	Fortran or advanced topics	22
3.3	<i>Project-based teaching</i>	23

## II C++ 25

4	<b>Basic elements of C++</b>	27
4.1	<i>Statements</i>	27
4.2	<i>Variables</i>	28
4.2.1	Variable declarations	28
4.2.2	Datatypes	29
4.2.3	Assignments	29
4.2.4	Floating point variables	30
4.2.5	Boolean values	30
4.3	<i>Input/Output, or I/O as we say</i>	30
4.4	<i>Expressions</i>	31
4.4.1	Truth values	32
4.4.2	Type conversions	32

<b>5</b>	<b>Conditionals</b>	35
5.1	<i>Conditionals</i>	35
5.1.1	<i>Switch statement</i>	36
5.1.2	<i>Scopes</i>	36
<b>6</b>	<b>Looping</b>	37
6.1	<i>Basic ‘for’ statement</i>	37
6.2	<i>Looping until</i>	38
<b>7</b>	<b>Functions</b>	41
7.1	<i>Parameter passing</i>	44
7.2	<i>Recursive functions</i>	46
7.3	<i>More function topics</i>	47
7.3.1	<i>Default arguments</i>	47
7.3.2	<i>Polymorphic functions</i>	48
<b>8</b>	<b>Scope</b>	49
8.1	<i>Lexical scope and dynamic scope</i>	49
8.2	<i>Static variables</i>	49
<b>9</b>	<b>Structures</b>	51
9.1	<i>Why structures?</i>	51
9.2	<i>The basics of structures</i>	51
<b>10</b>	<b>Classes and objects</b>	53
10.1	<i>What is an object?</i>	53
10.1.1	<i>Default constructor</i>	56
10.1.2	<i>Accessors</i>	56
10.1.3	<i>Methods</i>	57
10.2	<i>Inclusion relations between classes</i>	57
10.3	<i>Inheritance</i>	58
10.3.1	<i>Methods of base and derived classes</i>	59
<b>11</b>	<b>Arrays</b>	61
11.1	<i>Traditional arrays</i>	61
11.2	<i>Multi-dimensional arrays</i>	62
11.3	<i>Other allocation mechanisms</i>	62
11.4	<i>Vector class for arrays</i>	64
11.4.1	<i>Vector methods</i>	65
11.4.2	<i>Vectors are dynamic</i>	65
11.4.3	<i>Vector assignment</i>	65
11.4.4	<i>Vectors and functions</i>	66
11.4.5	<i>Dynamic size of vector</i>	66
11.4.6	<i>Timing</i>	67
11.5	<i>Wrapping a vector in an object</i>	68
11.6	<i>Multi-dimensional cases</i>	68
11.6.1	<i>Matrix as vector of vectors</i>	68
11.6.2	<i>Matrix class based on vector</i>	68
11.7	<i>Exercises</i>	69
<b>12</b>	<b>Strings</b>	71
12.1	<i>Basic string stuff</i>	71
12.2	<i>Conversion</i>	72

13	<b>Input/output</b>	73
13.1	<i>Formatted output</i>	73
13.2	<i>Floating point output</i>	75
13.3	<i>Saving and restoring settings</i>	77
13.4	<i>File output</i>	77
13.4.1	Output your own classes	78
13.5	<i>Input streams</i>	78
14	<b>References and addresses</b>	79
14.1	<i>Reference</i>	79
14.2	<i>Reference to class members</i>	80
14.3	<i>rvalue references</i>	80
15	<b>Polymorphism</b>	81
15.1	<i>The basic idea</i>	81
16	<b>Memory</b>	83
16.1	<i>Memory and scope</i>	83
17	<b>Pointers</b>	85
17.1	<i>What is a pointer</i>	85
17.2	<i>Pointers and addresses, C style</i>	85
17.2.1	Arrays and pointers	87
17.2.2	Pointer arithmetic	88
17.2.3	Multi-dimensional arrays	88
17.2.4	Parameter passing	88
17.2.5	Allocation	89
17.2.6	Memory leaks	90
17.3	<i>Safer pointers in C++</i>	91
18	<b>Prototypes</b>	93
18.1	<i>Prototypes for functions</i>	93
18.1.1	Header files	94
18.1.2	C and C++ headers	95
18.2	<i>Global variables</i>	95
18.3	<i>Prototypes for class methods</i>	96
18.4	<i>Header files and templates</i>	96
18.5	<i>Namespaces and header files</i>	96
19	<b>Namespaces</b>	97
19.1	<i>Solving name conflicts</i>	97
19.1.1	Namespace header files	98
20	<b>Preprocessor</b>	101
20.1	<i>Textual substitution</i>	101
20.2	<i>Parametrized macros</i>	102
20.3	<i>Conditionals</i>	102
20.3.1	Check on a value	103
20.3.2	Check for macros	103
21	<b>Templates</b>	105
21.1	<i>Templating over non-types</i>	106
22	<b>Error handling</b>	107
22.1	<i>General discussion</i>	107

22.2	<i>Exception handling</i>	107
23	<b>Standard Template Library</b>	111
23.1	<i>Containers</i>	111
23.1.1	<i>Iterators</i>	111
23.2	<i>Complex numbers</i>	112
23.3	<i>About the ‘using’ keyword</i>	112
24	<b>Obscure stuff</b>	113
24.1	<i>Casts</i>	113
24.1.1	<i>Casting constants</i>	113
24.1.2	<i>Dynamic cast</i>	113
24.1.3	<i>Legacy mechanism</i>	114
24.2	<i>Ivalue vs rvalue</i>	114
24.2.1	<i>Conversion</i>	115
24.2.2	<i>References</i>	115
24.2.3	<i>Rvalue references</i>	116
25	<b>More exercises</b>	117
25.1	<i>Practice</i>	117
25.2	<i>cplusplus</i>	117
25.3	<i>world best learning center</i>	117

### III Fortran 119

26	<b>Basics of Fortran</b>	121
26.1	<i>Main program</i>	121
26.1.1	<i>Program structure</i>	121
26.1.2	<i>Statements</i>	122
26.1.3	<i>Comments</i>	122
26.2	<i>Variables</i>	122
26.2.1	<i>Data types</i>	122
26.2.2	<i>Constants</i>	123
26.2.3	<i>Initialization</i>	123
27	<b>Loop constructs</b>	125
27.1	<i>Loop types</i>	125
27.2	<i>Interruptions of the control flow</i>	125
27.3	<i>Implied do-loops</i>	126
28	<b>Scope, subprograms, modules</b>	127
28.1	<i>Scope</i>	127
28.2	<i>Procedures</i>	127
28.2.1	<i>Subroutines and functions</i>	127
28.2.2	<i>Return results</i>	128
28.2.3	<i>Types of procedures</i>	128
28.2.4	<i>Optional arguments</i>	129
28.3	<i>Modules</i>	129
29	<b>Structures, eh, types</b>	131
30	<b>Classes and objects</b>	133
30.1	<i>Classes</i>	133

31	<b>Arrays</b>	139
31.1	<i>Static arrays</i>	139
31.1.1	Multi-dimensional	139
31.1.2	Query the size of an array	140
31.1.3	Allocatable arrays	140
31.2	Arrays to subroutines	140
31.3	Array slicing	141
31.4	Operating on an array	141
31.4.1	Arithmetic operations	141
31.4.2	Intrinsic functions	141
31.4.3	Restricting with <code>where</code>	142
32	<b>Pointers</b>	143
32.1	<i>Basic pointer operations</i>	143
32.2	<i>Example: linked lists</i>	144
33	<b>Input/output</b>	149
33.1	<i>Print to terminal</i>	149
33.1.1	Printing arrays	149
33.1.2	Formats	149
33.2	<i>File I/O</i>	150
34	<b>Array operations</b>	151
34.1	<i>Loops without looping</i>	151
34.1.1	Loops without dependencies	151
34.1.2	Loops with dependencies	152
IV	<b>Exercises and projects</b>	155
35	<b>Simple exercises</b>	157
35.1	<i>Looping exercises</i>	157
35.1.1	Further practice	158
35.2	<i>Object oriented exercises</i>	159
36	<b>Not-so simple exercises</b>	161
36.1	<i>List access</i>	161
37	<b>Prime numbers</b>	163
37.1	<i>Arithmetic</i>	163
37.2	<i>Conditionals</i>	163
37.3	<i>Looping</i>	163
37.4	<i>Functions</i>	164
37.5	<i>While loops</i>	164
37.6	<i>Global variables: optional</i>	164
37.7	<i>Structures</i>	165
37.8	<i>Classes and objects</i>	165
37.9	<i>Arrays</i>	166
38	<b>Infectuous disease simulation</b>	167
38.1	<i>Basic design</i>	167
39	<b>Geometry</b>	169
39.1	<i>Point class</i>	169

## Contents

---

39.2	<i>Using one class in another</i>	169
39.3	<i>Is-a relationship</i>	171
40	<b>PageRank</b>	173
40.1	<i>Basic ideas</i>	173
V	<b>Advanced topics</b>	175
41	<b>Tiniest of introductions to algorithms and data structures</b>	177
41.1	<i>Data structures</i>	177
41.1.1	Linked lists	177
41.1.2	Trees	178
41.2	<i>Algorithms</i>	180
41.2.1	Sorting	180
41.3	<i>Programming techniques</i>	181
41.3.1	Memoization	181
42	<b>Programming strategies</b>	183
42.1	<i>Programming: top-down versus bottom up</i>	183
42.1.1	Worked out example	184
42.2	<i>Coding style</i>	185
42.3	<i>Documentation</i>	185
42.4	<i>Testing</i>	185
43	<b>Complexity</b>	187
43.1	<i>Order of complexity</i>	187
43.1.1	Time complexity	187
43.1.2	Space complexity	187
44	<b>C for C++ programmers</b>	189
VI	<b>Index and such</b>	191
45	<b>Index</b>	193

## **PART I**

### **INTRODUCTION**



# **Chapter 1**

## **Introduction**

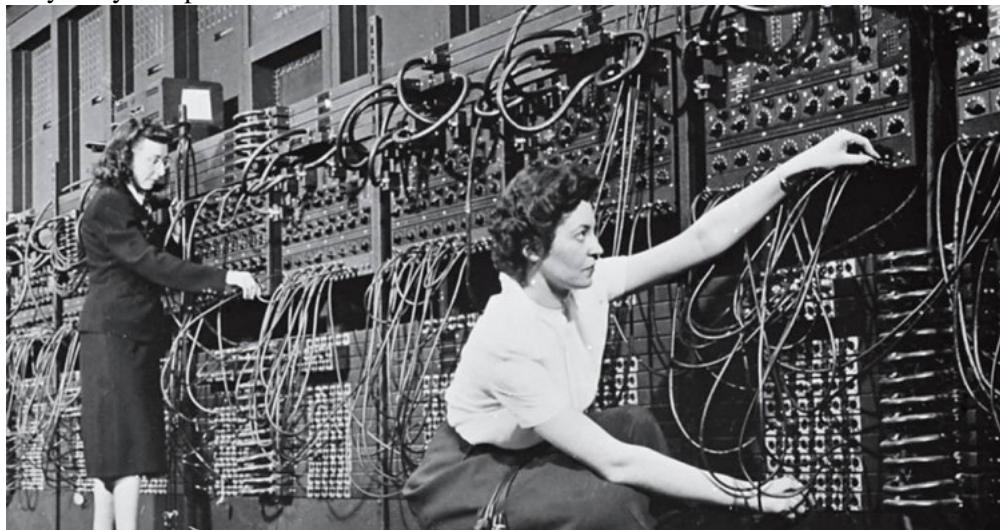
### **1.1 Programming and computational thinking**

#### **1.1.1 History**

Historically, computers were used for big physics calculations, for instance, atom bomb calculations



Very early computers were hardwired



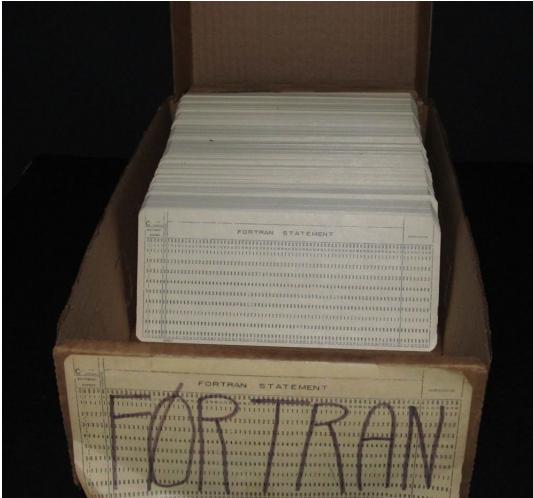
Later programs were written on punchcards

## 1. Introduction

---



Initial programming was about translating the math formulas; after a while they made a language for that: FORmula TRANslator



Programming is used in many different ways these days.

- You can make your own commands in Microsoft Word.
- You can make apps for your *smartphone*.
- You can solve the riddles of the universe using big computers.

This course is aimed at people in the last category.

### 1.1.2 Computational thinking

Programs can get pretty big



It's not just translating formulas anymore.

Translating ideas to computer code: computational thinking.

- Looking up a name in the phone book
  - start on page 1, then try page 2, et cetera
  - or start in the middle, continue with one of the halves.
- Elevator scheduling: someone at ground level presses the button, there are cars on floors 5 and 10; which one do you send down?
- The elevator programmer probably thinks: ‘if the button is pressed’, not ‘if the voltage on that wire is 5 Volt’.
- The Google car programmer probably writes: ‘if the car before me slows down’, not ‘if I see the image of the car growing’.
- ... but probably another programmer had to write that translation.

What is the structure of the data in your program?

Stack: you can only get at the top item



Queue: items get

added in the back, processed at the front



### 1.1.3 Hardware

Yes, it's there, but we don't think too much about it in this course.

<https://youtu.be/JEpsKnWZrJ8>

### 1.1.4 Algorithms

An algorithm is a sequence of unambiguous instructions for solving a problem, i.e., for obtaining a required output for any legitimate input in a finite amount of time [A. Levitin, Introduction to The Design and Analysis of Algorithms, Addison-Wesley, 2003]

The instructions are in some language:

- We will teach you C++ and Fortran;
- the compiler translates those languages to machine language
- Abstraction: a program often defines its own language that implements concepts of your application.
- Simple instructions: arithmetic.
- Complicated instructions: control structures
  - conditionals
  - loops
- Input and output data: to/from file, user input, screen output, graphics.
- Data during the program run:
  - Simple variables: character, integer, floating point
  - Arrays: indexed set of characters and such
  - Data structures: trees, queues
    - \* Defined by the user, specific for the application
    - \* Found in a library (big difference between C/C++)

## 1.2 About the choice of language

There are many programming languages, and not every language is equally suited for every purpose. In this book you will learn C++ and Fortran, because they are particularly good for scientific computing. And by 'good', we mean

- They can express the sorts of problems you want to tackle in scientific computing, and
- they execute your program efficiently.

There are other languages that may not be as convenient or efficient in expressing scientific problems. For instance, *python* is a popular language, but typically not the first choice if you're writing a scientific program. As an illustration, here is simple sorting algorithm, coded in both C++ and python.

Python vs C++ on bubblesort:

```
for i in range(n-1):
    for j in range(n-i-1):
        if numbers[j+1]<numbers[j]:
            swap = numbers[j+1]
            numbers[j+1] = numbers[j]
            numbers[j] = swap
for (int i=0; i<n-1; i++)
    for (int j=0; j<n-1-i; j++)
        if (numbers[j+1]<numbers[j]) {
            int swap = numbers[j+1];
            numbers[j+1] = numbers[j];
            numbers[j] = swap;
        }
[] python bubblesort.py 5000
Elapsed time: 12.1030311584
[] ./bubblesort 5000
Elapsed time: 0.24121
```

But this ignores one thing: the sorting algorithm we just implemented is not actually a terribly good one, and in fact python has a better one built-in.

Python with quicksort algorithm:

```
numpy.sort(numbers,kind='quicksort')

[] python arraysort.py 5000
Elapsed time: 0.00210881233215
```

So that is another consideration when choosing a language: is there a language that already comes with the tools you need.

### 1.3 Further reading

Tutorial, assignments: <http://www.cppforschool.com/>

Problems to practice: <http://www.spoj.com/problems/classical/>



# Chapter 2

## Warming up

### 2.1 Programming environment

Programming can be done in any number of ways. It is possible to use an Integrated Development Environment (IDE) such as *Xcode* or *Visual Studio*, but for the purposes of this course you should really learn a *Unix* variant.

- If you have a *Linux* computer, you are all set.
- If you have an *Apple* computer, it is easy to get you going. Install *XQuartz* and a *package manager* such as *homebrew* or *macports*.
- *Microsoft Windows* users can use *putty* but it is probably a better solution to install a virtual environment such as *VMware* (<http://www.vmware.com/>) or *Virtualbox* (<https://www.virtualbox.org/>).

Next, you should know a text editor. The two most common ones are *vi* and *emacs*.

#### 2.1.1 Language support in your editor

The author of this book is very much in favour of the *emacs* editor. The main reason is its support for programming languages. Most of the time it will detect what language a file is written in, based on the file extension:

- `cxx`, `cpp`, `cc` for C++, and
- `f90`, `F90` for Fortran.

If your editor somehow doesn't detect the language, you can add a line at the top of the file:

```
// -*- c++ -*-
```

for C++ mode, and

```
! -*- f90 -*-
```

for Fortran mode.

Main advantages are automatic indentation (C++ and Fortran) and supplying block end statements (Fortran). The editor will also apply ‘syntax colouring’ to indicate the difference between keywords and variables.

## 2.2 Compiling

The word ‘program’ is ambiguous. Part of the time it means the *source code*: the text that you type, using a text editor. And part of the time it means the *executable*, a totally unreadable version of your source code, that can be understood and executed by the computer. The process of turning your source code into an executable is called *compiling*, and it requires something called a *compiler*. (So who wrote the source code for the compiler? Good question.)

### 2.2.1 Compiling C++

Let’s say that

- you have a source code file `myprogram.cxx`,
- and you want an executable file called `myprogram`,
- and your compiler is `g++`, the C++ compiler of the *GNU* project. (If you have the Intel compilers, you will use `icpc` instead.)

To compile your program, you then type

```
g++ -o myprogram myprogram.cxx
```

On TACC machines, use the Intel compiler:

```
icpc -o myprogram myprogram.cxx
```

which you can verbalize as ‘invoke `g++`, with output `myprogram`, on `myprogram.cxx`’.

So let’s do an example.

This is a minimal program:

```
#include <iostream>
using namespace std;

int main() {
    return 0;
}
```

1. The first two lines are magic, for now. Always include them.
2. The `main` line indicates where the program starts; between its opening and closing brace will be the *program statements*.
3. The `return` statement indicates successful completion of your program.

As you may have guessed, this program does absolutely nothing.

Here is a statement that at least produces some output:

```
cout << "Hello world!" << endl;
```

**Exercise 2.1.** Make a program source file that contains the ‘hello world’ statement, compile it and run it. Think about where the statement goes.

## 2.2.2 Compiling Fortran

For Fortran programs, the compiler is `gfortran` for the GNU compiler, and `ifort` for Intel. Fortran programs can have a number of extensions, but some of them have special meaning to the compiler, by convention. In this book we adopt the `F90` extension.

The minimal Fortran program is:

```
Program SomeProgram
    ! stuff goes here
End Program SomeProgram
```

**Exercise 2.2.** Add the line

```
print *, "Hello world!"
```

to the empty program, and compile and run it.

## 2. Warming up

## **Chapter 3**

### **Teachers guide**

This book was written for a one-semester introductory programming course at The University of Texas at Austin, aimed primarily at students in the physical and engineering sciences. Thus, examples and exercises are as much as possible scientifically motivated. This target audience also explains the inclusion of Fortran.

This book is not encyclopedic. Rather than teaching each topic in its full glory, the author has taken a ‘good practices’ approach, where students learn enough of each topic to become a competent programmer. This serves to keep this book at a manageable length, and to minimize class lecture time, emphasizing lab exercises instead.

Even then, there is more material here than can be covered and practiced in one semester. If only C++ is taught, it is probably possible to cover the whole of Part II; for the case where both C++ and Fortran are taught, we have a suggested timeline below.

#### **3.1 Justification**

The chapters of Part II and Part III are presented in suggested teaching order. Here we briefly justify our (non-standard) sequencing of topics and outline a timetable for material to be covered in one semester. Most notably, Object-Oriented programming is covered before arrays and pointers come very late, if at all.

There are several thoughts behind this. For one, dynamic arrays in C++ are most easily realized through the `std::vector` mechanism, which requires an understanding of classes. The same goes for `std::string`.

Secondly, in the traditional approach, OOP is taught late, if at all, in the course. We consider OOP to be an important notion in program design, and central to C++, rather than an embellishment on the traditional C mechanisms.

#### **3.2 Timeline for a C++/F03 course**

As remarked above, this book is based on a course that teaches both C++ and Fortran2003. Here we give the timeline used, including some of the assigned exercises.

### 3. Teachers guide

---

lesson#	date	Topic	Exercises	prime	geom	infect
1	1/18	Statements and expressions	<a href="#">37.1</a> , <a href="#">37.2</a>			
2	1/24	Conditionals	<a href="#">37.3</a>			
3	1/26	Control structures				
4	1/31	Looping	<a href="#">37.4</a> , <a href="#">37.5</a> , <a href="#">37.7</a>			
5		continue				
6	2/06	Functions	<a href="#">37.6</a>			
7		continue				
8	2/12	I/O (lecture 8)				
9	2/19	Structs	<a href="#">37.9</a>			
10	2/23	Objects	<a href="#">37.10</a> , <a href="#">37.12</a>		<a href="#">39.1</a>	
11		continue				
12	2/28	has-a relation			<a href="#">39.4</a> , <a href="#">39.5</a> , <a href="#">39.1</a> , <a href="#">39.6</a>	
13	3/02	inheritance			<a href="#">39.7</a> , <a href="#">39.8</a>	
14	3/07	Arrays				<a href="#">38.1</a> , <a href="#">38.2</a> , <a href="#">38.3</a> , <a href="#">38.4</a>
15		continue				
16	3/23	Strings				

Table 3.1: Two-month lesson plan for C++

For a one semester course of slightly over three months, two months would be spent on C++ (see table 3.1), after which a month is enough to explain Fortran. Remaining time will go to exams and elective topics.

We also give three sets of exercises that form a sequence of assignments that build on each other:

**prime** Prime number testing, culminating in prime number sequence objects, and testing a corollary of the Goldbach conjecture.

**geom** Geometry related concepts; this is mostly an exercise in object-oriented programming.

**infect** The spreading of infectious diseases; these are exercises in basic array programming.

#### 3.2.1 Fortran or advanced topics

After two months of grounding in OOP programming in C++, the Fortran lectures and exercises reprise this sequence, letting the students do the same exercises in Fortran that they did in C++. However, array mechanisms in Fortran warrant a separate lecture.

If the course focuses solely on C++, the third month can be devoted to

- templates,
- exceptions,
- namespaces,
- multiple inheritance,
- the cpp preprocessor,
- closures.

### 3.3 Project-based teaching

To an extent it is inevitable that students will do a number of exercises that are not connected to any earlier or later ones. However, to give some continuity, we have given some programming projects that students gradually build towards.

Currently most successful is the ‘prime numbers’ sequence; chapter 37.

Rather than including the project exercises in the didactic sections, each section in these projects list the prerequisite basic sections.

### 3. Teachers guide

## **PART II**

**C++**



# Chapter 4

## Basic elements of C++

### 4.1 Statements

Each programming language has its own (very precise!) rules for what can go in a source file. Globally we can say that a program contains instructions for the computer to execute, and these instructions take the form of a bunch of ‘statements’. Here are some of the rules on statements; you will learn them in more detail as you go through this book.

- A program contains statements, each terminated by a semicolon.
- ‘Curly braces’ can enclose multiple statements.
- A statement can correspond to some action when the program is executed.
- Some statements are definitions, of data or of possible actions.
- Comments are ‘Note to self’, short:

```
cout << "Hello world" << endl; // say hi!
```

and arbitrary:

```
cout << /* we are now going
           to say hello
        */ "Hello!" << /* with newline */ endl;
```

**Exercise 4.1.** Take the ‘hello world’ program you wrote earlier, and duplicate the hello-line.

Compile and run.

Does it make a difference whether you have the two hellos on the same line or on different lines?

Experiment with other changes to the layout of your source. Find at least one change that leads to a compiler error.

A program can not be anything: there are some fixed ingredients, and some rules governing the form of what you are saying.

You see that certain parts of your program are inviolable:

- There are *keywords* such as `return` or `cout`; you can not change their definition.
- Curly braces and parentheses need to be matched.
- There has to be a `main` keyword.
- The `iostream` and `namespace` are usually needed.

**Exercise 4.2.** Experiment with the `cout` statement. Replace the string by a number or a mathematical expression. Can you guess how to print more than one thing, for instance the string `One third is` and the result of  $1/3$ , with the same `cout` statement?

## 4.2 Variables

A program could not do much without storing data: input data, temporary data for intermediate results, and final results. Data is stored in *variables*, which have

- a name, so that can refer to them, and
- a *datatype*.

Think of a variable as a labeled place in memory.

- The variable is defined in a *variable declaration*,
- which can include an *variable initialization*.
- After a variable is defined, and given a value, it can be used,
- or given a (new) value in a *variable assignment*.

```
int i, j; // declaration
i = 5; // set a value
i = 6; // set a new value
j = i+1; // use the value of i
i = 7; // change the value of i
        // but this doesn't affect j
```

### 4.2.1 Variable declarations

A variable is defined once in a *variable declaration*, but it can be given a (new) value multiple times. It is not an error to use a variable that has not been given a value, but it may lead to strange behaviour at runtime.

Programs usually contain data, which is stored in a *variable*. A variable has

- a *datatype*,
- a name, and
- a value.

These are defined in a *variable declaration* and/or *variable assignment*.

- A variable name has to start with a letter,
- can contains letters and digits, but not most special characters (except for the underscore).
- For letters it matters whether you use upper or lowercase: the language is *case sensitive*.

There are a couple of ways to make the connection between a name and a type. Here is a simple *variable declaration*, which establishes the name and the type:

```
int n;
float x;
int n1,n2;
double re_part,im_part;
```

Declarations can go pretty much anywhere in your program, but need to come before use of the variable.

### 4.2.2 Datatypes

Variables come in different types;

- We call a variable of type `int`, `float`, `double` a *numerical variable*.
- For characters: `char`. Strings are complicated.
- You can make your own types. Later.

### 4.2.3 Assignments

Setting a variable

```
i = 5;
```

means storing a value in the memory location. It is not defining a mathematical equality

```
let i = 5.
```

Once you have declared a variable, you need to establish a value. This is done in an *assignment statement*. After the above declarations, the following are legitimate assignments:

```
n = 3;
x = 1.5;
n1 = 7; n2 = n1 * 3;
```

You see that you can assign both a simple value or an *expression*; see section 4.4 for more detail.

A variable can be given a value more than once. You the following sequence of statements is a legitimate part of a program:

```
int n;
n = 3;
n = 2*n + 5;
n = 3*n + 7;
```

Update:

```
x = x+2; y = y/3;
// can be written as
x += 2; y /= 3;
```

Integer add/subtract one:

```
i++; j--; /* same as: */ i=i+1; j=j-1;
```

Pre/post increment:

```
x = a[i++]; /* is */ x = a[i]; i++;
y = b[++i]; /* is */ i++; y = b[i];
```

You can also give a variable a value a in *variable initialization*. Confusingly, there are several ways of doing that. Here's two:

```
int n = 0;
double x = 5.3, y = 6.7;
double pi{3.14};
```

**Exercise 4.3.** Write a program that has several variables. Assign values either in an initialization or in an assignment. Print out the values.

#### 4.2.4 Floating point variables

Mathematically, integers are a special case of real numbers. In a computer, integers are stored very differently from *floating point* numbers.

- Within a certain range, roughly  $-2 \cdot 10^9, \dots, 2 \cdot 10^9$ , all integer values can be represented.
- On the other hand, not all real numbers have a floating point representation. For instance, since computer numbers are binary,  $1/2$  is representable but  $1/3$  is not.
- You can assign variables of one type to another, but this may lead to truncation (assigning a floating point number to an integer) or unexpected bits (assigning a single precision floating point number to a double precision).
- Default: `double`
- Float: `3.14f` or `3.14F`
- Long double: `1.411` or `1.41L`.

This prevents numerical accidents:

```
double x = 3.;
```

converts float to double, maybe introducing random bits.

#### 4.2.5 Boolean values

So far you have seen integer and real variables. There are also *boolean values* which represent truth values. There are only two values: `true` and `false`.

```
bool found{false};
```

**Exercise 4.4.** Print out `true` and `false`. What do you get?

### 4.3 Input/Output, or I/O as we say

A program typically produces output. For now we will only display output on the screen, but output to file is possible too. Regarding input, sometimes a program has all information for its computations, but it is also possible to base the computation on user input.

You have already seen `cout`:

```
float x = 5;
cout << "Here is the root: " << sqrt(x) << endl;
```

There is also a `cin`, which serves to take user input and put it in a numerical variable.

```
int i;
cin >> i;
```

However, this function is somewhat tricky. <http://www.cplusplus.com/forum/articles/6046/>.

It is better to use `getline`. This returns a string, rather than a value, so you need to convert it with the following bit of magic:

```
#include <iostream>
#include <sstream>
using namespace std;
/* ... */
std::string saymany;
int howmany;

cout << "How many times? ";
getline( cin, saymany );
stringstream saidmany(saymany);
saidmany >> howmany;
```

You can not use `cin` and `getline` in the same program.

**Exercise 4.5.** Write a program that

- Displays the message `Type a number,`
- accepts an integer number from you (use `cin`),
- and then prints out three times that number plus one.

For more I/O, see chapter 13.

## 4.4 Expressions

The most basic step in computing is to form expressions such as sums, products, logical conjunctions, string appending. Expressions in programming languages for the most part look the way you would expect them to.

- Mathematical operators: `+` `-` `/` and `*` for multiplication.
- C++ does not have a power operator (Fortran does).
- Integer modulus: `5%2`
- You can use parentheses: `5*(x+y)`. Use parentheses if you're not sure about the precedence rules for operators.
- ‘Power’ and various mathematical functions are realized through library calls.

Library calls:

```
#include <cmath>
```

```
x = pow(3, .5);
```

For squaring, usually better to write `x*x` than `pow(x, 2)`.

- Expression looks pretty much like in math.  
With integers:  $2+3$   
with reals:  $3.2/7$
- Use parentheses to group  $25.1 * (37+42/3.)$
- Careful with types.
- There is no ‘power’ operator: library functions. Needs a line

```
#include <cmath>
```

- Modulus: `%`

#### 4.4.1 Truth values

In addition to numerical types, there are truth values, `true` and `false`, with all the usual logical operators defined on them.

Logical expressions in C++ are evaluated using *shortcut operators*: you can write

```
x>=0 && sqrt(x)<2
```

If `x` is negative, the second part will never be evaluated because the ‘and’ conjunction can already be concluded to be false. Similarly, ‘or’ conjunctions will only be evaluated until the first true clause.

- Testing: `== != < > <= >=`
- Not, and, or: `!` `&&` `||`
- Shortcut operators:  
`if ( x>=0 && sqrt(x)<5 ) { }`

The ‘true’ and ‘false’ constants could strictly speaking be stored in a single bit. Normally, C++ does not do that, but there are bit operators that you can apply to, for instance, all the bits in an integer.

Bitwise: `&` `|` `^`

#### 4.4.2 Type conversions

Since a variable has one type, and will always be of that type, you may wonder what happens with

```
float x = 1.5;
int i;
i = x;
```

or

---

```
int i = 6;
float x;
x = i;
```

- Assigning a floating point value to an integer truncates the latter.
- Assigning an integer to a floating point variable fills it up with zeros after the decimal point.

**Exercise 4.6.**

- What happens when you assign a negative integer to a floating point variable?
- What happens when you assign a `float` to a `double`? Try various number for the original float. Can you explain the result? (Hint: think about the conversion between binary and decimal.)

The rules for type conversion in expressions are not entirely logical. Consider

```
float x; int i=5, j=2;
x = i/j;
```

This will give 2 and not 2.5, because `i / j` is an integer expression and is therefore completely evaluated as such, giving 2 after truncation. The fact that it is ultimately assigned to a floating point variable does not cause it to be evaluated as a computation on floats.

You can force the expression to be computed in floating point numbers by writing

```
x = (1.*i)/j;
```

or any other mechanism that forces a conversion, without changing the result. Another mechanism is the `cast`; this will be discussed in section 24.1.

Real to integer: round down:

```
double x,y; x = ....; y = ....;
int i; i = x+y;
```

Dangerous:

```
int i,j; i = ...; j = ...;
double x; x = 1+i/j;
```

The fraction is executed as integer division.

For floating point result do:

```
(double)i/j /* or */ (1.*i)/j
```

**Exercise 4.7.** Write two programs, one that reads a temperature in Centigrade and converts to Fahrenheit, and one that does the opposite conversion.

$$C = (F - 32) \cdot 5/9, \quad F = 9/5 C + 32$$

Can you use Unix pipes to make one accept the output of the other?

**Exercise 4.8.** Write a program that ask for two integer numbers `n1, n2`.

- Assign the integer ratio  $n_1/n_2$  to a variable.
- Can you use this variable to compute the modulus

$n_1 \mod n_2$

(without using the % modulus operator!)

Print out the value you get.

- Also print out the result from using the modulus operator: %.

Complex numbers exist, see section [23.2](#).

# Chapter 5

## Conditionals

### 5.1 Conditionals

A program consisting of just a list of assignment and expressions would not be terribly versatile. At least you want to be able to say ‘if  $x > 0$ , compute, otherwise compute something else’, or ‘until some test is true, iterate the following computations’. The mechanism for testing and choosing an action accordingly is called a *conditional*.

Here are some forms. Single statement

```
if (x<0)
    x = -x;
```

Single statement and alternative:

```
if (x>=0)
    x = 1;
else
    x = -1;
```

Multiple statements:

```
if (x<0) {
    x = 2*x; y = y/2;
} else {
    x = 3*x; y = y/3;
}
```

Chaining conditionals (where the dots stand for omitted code):

```
if (x>0) {
    ....
} else if (x<0) {
    ....
} else {
    ....
}
```

## 5. Conditionals

---

Nested conditionals:

```
if (x>0) {  
    if (y>0) {  
        ....  
    } else {  
        ....  
    }  
} else {  
    ....  
}
```

(In that last example the outer curly brackets are optional. But it's safer to use them anyway.)

### 5.1.1 Switch statement

If you have a number of cases corresponding to specific integer values, there is the `switch` statement.

```
switch (n) {  
case 1 :  
case 2 : cout << "very small" << endl;  
break;  
case 3 : cout << "trinity" << endl;  
break;  
default : cout << "large" << endl;  
}
```

### 5.1.2 Scopes

The true and false branch of a conditional can consist of a single statement, or of a block in curly brackets. Such a block creates a *scope*! in conditional branches where you can define local variables.

```
if ( something ) {  
    int i;  
    .... do something with i  
}  
// the variable 'i' has gone away.
```

**Exercise 5.1.** Read in an integer. If it's a multiple of three print 'Fizz'; if it's a multiple of five print 'Buzz'. If it is a multiple of both three and five print 'FizzBuzz'. Otherwise print nothing.

# Chapter 6

## Looping

### 6.1 Basic ‘for’ statement

There are many cases where you want to repeat an operation or series of operations:

- A time-dependent numerical simulation executes a fixed number of steps, or until some stopping test.
- Recurrences:

$$x_{i+1} = f(x_i).$$

- Inspect or transformation every element of a database table.

(Fine point: the first two cases actually need to be performed in sequences, while the last one corresponds more to a mathematical ‘forall’ quantor.)

What you need is known as a *loop*: a number of statements that get repeated. The two types of loop statement in C++ are:

- the *for loop* which is typically associated with a pre-determined number of repetitions, and with the repeated statements being based on a counter of some sort; and
- the *while loop*, where the statements are repeated indefinitely until some condition is satisfied.

We will now consider the for loop; the while loop comes in section 6.2.

In the most common case, a for loop has a *loop counter*, ranging from some lower to some upper bound. And example showing the syntax for this simple case is:

```
for (int var=low; var<upper; var++) {  
    // statements involving the loop variable:  
    cout << "The square of " << var << " is " << var*var << endl;  
}
```

The `for` line is called the *loop header*, and the statements between curly brackets the *loop body*.

**Exercise 6.1.** Read an integer value, and print ‘Hello world’ that many times.

The loop header has three components, all of which are optional.

- An initialization. This is usually a declaration and initialization of an integer *loop variable*. Using floating point values is less advisable.
- A stopping test, usually involving the loop variable. If you leave this out, you need a different mechanism for stopping the loop; see 6.2.

## 6. Looping

---

- An increment. You can let the loop count down by using `i--`.

Some variations on the simple case.

- The loop variable can be defined outside the loop:

```
int var;  
for (var=low; var<upper; var++) {
```

but it's cleaner to make it local. However:

```
int var;  
..... code that sets var .....\nfor ( ; var<upper; var++) {  
    ...  
}
```

- The stopping test can be omitted

```
for (int var=low; ; var++) { ... }
```

if the loops ends in some other way. You'll see this later.

- The stopping test doesn't need to be an upper bound:

```
for (int var=high-1; var>=low; var--) { ... }
```

- Here are a couple of popular increments:

- `i++` for a loop that counts forward;
- `i--` for a loop that counts backward;
- `i+=2` to cover only odd or even numbers, depending on where you started;
- `i*=10` to cover powers of ten.

Quite often, the loop body will contain another loop. For instance, you may want to iterate over all elements in a matrix. Both loops will have their own unique loop variable.

```
for (int i=0; i<m; i++)  
    for (int j=0; j<n; j++)  
        ...
```

**Exercise 6.2.** Write an  $i, j$  loop that prints out all pairs with

$$1 \leq i \leq 10, \quad 1 \leq j < i.$$

Same, but

$$1 \leq i \leq 10, |i - j| < 2.$$

### 6.2 Looping until

The basic for loop looks pretty deterministic: a loop variable ranges through a more-or-less prescribes set of values. This is appropriate for looping over the elements of an array, but not if you are coding some process that needs to run until some dynamically determined condition is satisfied. In this section you will see some ways of coding such cases.

First of all, the stopping test in the ‘for’ loop is optional:

```
for (int var=low; ; var=var+1) { ... }
```

So how do you end such a loop? For that you use the `break` statement. If the execution encounters this statement, it will continue with the first statement after the loop.

```
for (int var=low; ; var=var+1) {
    // statements;
    if (some_test) break;
    // more statements;
}
```

**Exercise 6.3.** Write a double loop over  $0 \leq i, j < 10$  that prints the first pair where the product of indices satisfies  $i \cdot j > 80$ .

Another mechanism to alter the control flow in a loop is the `continue` statement. If this is encountered, execution skips to the start of the next iteration.

```
for (int var=low; var<N; var++) {
    statement;
    if (some_test) {
        statement;
        statement;
    }
}
```

Alternative:

```
for (int var=low; var<N; var++) {
    statement;
    if (!some_test) continue;
    statement;
    statement;
}
```

The other possibility is a `while` loop, which repeats until a condition is met.

Syntax:

```
while ( condition ) {
    statements;
}
```

or

```
do {
    statements;
} while ( condition );
```

The `while` loop does not have a counter or an update statement; if you need those, you have to create them yourself.

## 6. Looping

---

The two while loop variants can be described as ‘pre-test’ and ‘post-test’. The choice between them entirely depends on context. Here is an example in which the second syntax is more appropriate.

```
cout << "Enter a positive number: " ;
cin >> invar;
while (invar>0) {
    cout << "Enter a positive number: " ;
    cin >> invar;
}
cout << "Sorry, " << invar << " is negative" << endl;
```

Problem: code duplication.

```
do {
    cout << "Enter a positive number: " ;
    cin >> invar;
} while (invar>0);
cout << "Sorry, " << invar << " is negative" << endl;
```

More elegant.

Exercise 6.4. One bank account has 100 dollars and earns a 5 percent per year interest rate.

Another account has 200 dollars but earns only 2 percent per year. In both cases the interest is deposited into the account.

After how many years will the amount of money in both accounts be the same?

## Chapter 7

### Functions

A *function* (or *subprogram*) is a way to abbreviate a block of code and replace it by a single line.

- Find a block of code that has a clearly identifiable function.
- Turn this into a function: the function definition will contain that block, plus a header and (maybe) a return statement.
- The function definition is placed before the main program.
- The function is called by its name.

Example: the code for an odd/even test

```
for (int i=0; i<N; i++) {
    cout << i;
    if (i%2==0)
        cout << " is even";
    else
        cout << " is odd";
    cout << endl;
}
```

becomes

```
void report_evenness(int n) {
    cout << i;
    if (i%2==0)
        cout << " is even";
    else
        cout << " is odd";
    cout << endl;
}
...
int main() {
    ...
    for (int i=0; i<N; i++)
        report_evenness(i);
}
```

In the *function definition*:

- The keyword `void` indicates that the function does not give any results back to the main program.
- The name `report_evenness` is picked by you.
- The parenthetical `(int n)` is called the ‘argument list’ or ‘parameter list’: it says that the function takes an `int` as input. For purposes of the function, the `int` will have the name `n`, but this is not necessarily the same as the name in the main program.
- The ‘body’ of the function, the code that is going to be executed, is enclosed in curly brackets.

The *function call* consists of

- The name of the function, and
- In between parentheses, the value of the input argument.

In the previous example, the function had an input, and performed some screen output. To have a function that takes part in the computation of your program, you would write something like: Functions are defined before the main program, and used in that program: Here is a program with a function that doubles its input:

```
#include <iostream>
using namespace std;

int twice_function(int n) {
    int twice_the_input = 2*n;
    return twice_the_input;
}

int main() {
    int number = 3;
    cout << "Twice three is: " <<
        twice_function(number) << endl;
    return 0;
}
```

Functions can be motivated as making your code more structured and intelligible. The source where you use the function call becomes shorter, and the function name makes the code more descriptive. This is sometimes called ‘self-documenting code’.

Sometimes introducing a function can be motivated from a point of *code reuse*: if the same block of code appears in two places in your source, you replace this by one function definition, and two (single line) function calls. The two occurrences of the function code do not have to be identical:

```
double x, y, v, w;
y = ..... computation from x .....
w = ..... same computation, but from v .....
```

would be replaced by

```
double computation(double in) {
    return .... computation from 'in' ....
}
```

---

```
y = computation(x);
w = computation(v);
```

A final argument for using functions is code maintainability:

- Easier to debug: if you use the same (or roughly the same) block of code twice, and you find an error, you need to fix it twice.
- Maintenance: if a block occurs twice, and you change something in such a block once, you have to remember to change the other occurrences too.
- Localization: any variables that only serve the calculation in the function now have a limited scope.

```
void print_mod(int n,int d) {
    int m = n%d;
    cout << "The modulus of " << n << " and " << d
        << " is " << m << endl;
```

Loosely, a function takes input and computes some result which is then returned. Formally, a function consists of:

- *function result type*: you need to indicate the type of the result;
- name: you get to make this up;
- zero or more *function parameters* or *function arguments*: the data that the function operates on; and the
- *function body*: the statements that make up the function. The function body is a *scope*: it can have local variables. (You can not nest function definitions.)
- a *return statement*. Which doesn't have to be the last statement, by the way.

The function can then be used in the main program, or in another function:

The function call

1. causes the function body to be executed, and
2. the function call is replaced by whatever you *return*.
3. (If the function does not return anything, for instance because it only prints output, you declare the return type to be `void`.)

```
void print_header() {
    cout << "*****" << endl;
    cout << "* Ouput      *" << endl;
    cout << "*****" << endl;
}
int main() {
    print_header();
    cout << "The results for day 25:" << endl;
    // code that prints results ....
    return 0;
}

void print_header(int day) {
    cout << "*****" << endl;
```

```
cout << "* Ouput      *" << endl;
cout << "*****" << endl;
cout << "The results for day " << day << ":" << endl;
}
int main() {
    print_header(25);
    // code that prints results ....
    return 0;
}

#include <cmath>
double pi() {
    return 4*atan(1.0);
}
```

A function body defines a *scope*: the local variables of the function calculation are invisible to the calling program.

Functions can not be nested: you can not define a function inside the body of another function.

## 7.1 Parameter passing

C++ functions resemble mathematical functions: you have seen that a function can have an input and an output. In fact, they can have multiple inputs. The following style of programming is very much inspired by mathematical functions, and is known as *functional programming*:

- A function has one result, which is returned through a return statement. The function call then looks like

```
y = f(x1, x2, x3);
```

- The definition of the C++ parameter passing mechanism says that input arguments are copied to the function, meaning that they don't change in the calling program:

```
x = 5;
y = f(x);
// x is still 5
```

We say that input argument as *passed by value* *pass by value*.

**Exercise 7.1.** Early computers had no hardware for computing a square root. Instead, they used *Newton's method*. Suppose you want to compute

$$x = \sqrt{y}.$$

This is equivalent to finding the zero of

$$f(x) = x^2 - y.$$

Newton's method does this by evaluating

$$x_{\text{next}} = x - f(x)/f'(x)$$

until the guess is accurate enough.

- Write functions `f(x, y)` and `deriv(x, y)`, and a function `newton_root` that uses `f` and `deriv` to iterate to some precision.
- Take an initial guess for `x`, not zero.
- As a stopping test, use  $|f(x, y)| < 10^{-5}$ .
- Use `double` for all your variables.

Having only one output is a limitation on functions. Therefore there is a mechanism for altering the input parameters and returning (possibly multiple) results that way. To do this, you use an ampersand for the parameter in the function definition:

```
void f(int &i) {
    i = /* some expression */ ;
}
int main() {
    int i;
    f(i);
    // i now has the value that was set in the function
}
```

Using the ampersand, the parameter is *passed by reference*: instead of copying the value, the function receives a ‘reference’: information where the variable is stored in memory.

As an example, consider a function that tries to read a value from a file. With anything file-related, you always have to worry about the case of the file not existing and such. So our function return:

- a boolean value to indicate whether the read succeeded, and
- the actual value if the read succeeded.

The following is a common idiom, where the success value is returned through the `return` statement, and the value through a parameter.

```
bool can_read_value( int &value ) {
    int file_status = try_open_file();
    if (file_status==0)
        value = read_value_from_file();
    return file_status!=0;
}
...
if (!can_read_value(n))
    // if you can't read the value, set a default
    n = 10;
```

**Exercise 7.2.** Write a function `swap` of two parameters that exchanges the input values:

```
int i=2, j=3;
swap(i, j);
```

```
// now i==3 and j==2
```

**Exercise 7.3.** Write a function that tests divisibility and returns a remainder:

```
int number, divisor, remainder;
// get the number and divisor from the user
if ( is_divisible(number, divisor, remainder) )
    cout << number << " is divisible by " << divisor << endl;
else
    cout << number << "/" << divisor <<
        " has remainder " << remainder << endl;
```

## 7.2 Recursive functions

In mathematics, sequences are often recursively defined. For instance, the sequence of factorials  $n \mapsto f_n \equiv n!$  can be defined as

$$f_0 = 1, \quad \forall_{n>0}: f_n = n \times f_{n-1}.$$

Instead of using a subscript, we write an argument in parentheses

$$F(n) = n \times F(n - 1) \quad \text{if } n > 0, \text{ otherwise } 1$$

and now it looks like a C++ function:

```
int factorial(int n)
```

is the function header of a factorial function. So what would the function body be? We need a `return` statement, and what we return should be  $n \times F(n - 1)$ :

```
int factorial(int n) {
    return n * factorial(n-1);
} // almost correct, but not quite
```

So what happens if you write

```
int f3; f3 = factorial(3);
```

Well,

- The expression `factorial(3)` calls the `factorial` function, substituting 3 for the argument `n`.
- The return statement returns `n * factorial(n-1)`, in this case `3 * factorial(2)`.
- But what is `factorial(2)`? Evaluating that expression means that the `factorial` function is called again, but now with `n` equal to 2.
- Evaluating `factorial(2)` returns `2 * factorial(1), ...`
- ... which returns `1 * factorial(0), ...`
- ... which return ...
- Uh oh. We forgot to include the case where `n` is zero. Let's fix that:

```
int factorial(int n) {
    if (n==0)
        return 1;
    else
        return n*factorial(n-1);
}
```

- Now `factorial(0)` is 1, so `factorial(1)` is `1*factorial(0)`, which is 1,...
- ... so `factorial(2)` is 2, and `factorial(3)` is 6.

Exercise 7.4. The sum of squares:

$$S_n = \sum_{n=1}^N n^2$$

can be defined recursively as

$$S_1 = 1, \quad S_n = n^2 + S_{n-1}.$$

Write a recursive function that implements this second definition. Test it on numbers that are input by the user.

Then write a program that prints the first 100 sums of squares.

Exercise 7.5. Write a recursive function for computing Fibonacci numbers:

$$F_0 = 1, \quad F_1 = 1, \quad F_n = F_{n-1} + F_{n-2}$$

First write a program that computes  $F_n$  for a value  $n$  that is input by the user.

Then write a program that prints out a sequence of Fibonacci numbers; the user should input how many.

If you let your Fibonacci program print out each time it computes a value, you'll see that most values are computed several times. (Math question: how many times?) This is wasteful in running time. This problem is addressed in section 41.3.1.

## 7.3 More function topics

### 7.3.1 Default arguments

Functions can have *default argument(s)*:

```
double distance( double x, double y=0. ) {
    return sqrt( (x-y)*(x-y) );
}
```

Any default argument(s) should come last in the parameter list.

### 7.3.2 Polymorphic functions

You can have multiple functions with the same name:

```
double sum(double a, double b) {  
    return a+b; }  
double sum(double a, double b, double c) {  
    return a+b+c; }
```

Distinguished by input parameters: can not differ only in return type.

# Chapter 8

## Scope

### 8.1 Lexical scope and dynamic scope

### 8.2 Static variables

Variables in a function have *lexical scope* limited to that function. Normally they also have *dynamic scope* limited to the function execution: after the function finishes they completely disappear. (Class objects have their *destructor* called.)

There is an exception: a *static variable* persists between function invocations.

```
void fun() {
    static int remember;
}
```

For example

```
int onemore() {
    static int remember++; return remember;
}
int main() {
    for ( ... )
        cout << onemore() << endl;
    return 0;
}
```

gives a stream of integers.

Exercise 8.1. The static variable in the `onemore` function is never initialized. Can you find a mechanism for doing so? Can you do it with a default argument to the function?

## 8. Scope

---

# Chapter 9

## Structures

### 9.1 Why structures?

You have seen the basic datatypes in section 4.2.3. These are enough to program whatever you want, but it would be nice if the language had some datatypes that are more abstract, closer to the terms in which you think about your application. For instance, if you are programming something to do with geometry, you had rather talk about points than explicitly having to manipulate their coordinates.

Structures are a first way to define your own datatypes. A `struct` acts like a datatype for which you choose the name. A `struct` contains other datatypes; these can be elementary, or other structs.

```
struct vector { double x; double y; } ;
```

The elements of a structure are also called *members*/*member (of struct)*. You can give them an initial value:

```
struct vector { double x=0.; double y=0.; } ;
```

### 9.2 The basics of structures

Once you have defined a structure, you can make variables of that type. Setting and initializing them takes a new syntax:

```
struct vector p1,p2;  
  
p1.x = 1.; p1.y = 2.;  
p2 = {3.,4.};  
  
p2 = p1;
```

You can pass a structure to a function:

```
double distance( struct vector p1,struct vector p2 ) {  
    double d1 = p1.x-p2.x, d2 = p1.y-p2.y;  
    return sqrt( d1*d1 + d2*d2 );  
}
```

## 9. Structures

---

Prevent copying cost by passing by reference, use `const` to prevent changes:

```
double distance( const struct vector &p1,const struct vector &p2 ) {  
    double d1 = p1.x-p2.x, d2 = p1.y-p2.y;  
    return sqrt( d1*d1 + d2*d2 );  
}
```

You can return a structure from a function:

```
struct vector vector_add  
    ( struct vector p1,struct vector p2 ) {  
    struct vector p_add = {p1.x+p2.x,p1.y+p2.y};  
    return p_add;  
};
```

(Something weird here with scopes: the explanation is that the returned value is copied.)

**Exercise 9.1.** Write a function `inner_product` that takes two `vector` structures and computes the inner product.

**Exercise 9.2.** Write a  $2 \times 2$  matrix class (that is, a structure storing 4 real numbers), and write a function `multiply` that multiplies a matrix times a vector.

# Chapter 10

## Classes and objects

### 10.1 What is an object?

You learned about `structs` as a way of abstracting a little from the elementary data types. The elements of a structure were called its members.

You also saw that it is possible to write functions that work on structures. Since these functions are really tied to the definition of the `struct`, shouldn't there be a way to make that tie explicitly?

That's what an object is:

- An object is like a structure in that it has data members.
- An object has *methods* which are the functions that operate on that object.

C++ does not actually have a ‘object’ keyword; instead you define a class with the `class` which describes all the objects of that class.

First of all, you can make an object look pretty much like a structure:

```
class Vector {  
public:  
    double x, y;  
};  
  
int main() {  
    Vector p1;  
    p1.x = 1.; p1.y = 2.;
```

- There are data members. We will get to the `public` in a minute.
- You make an object of that class by using the class name as the datatype.
- The data members can be accessed with the period.

Next we'll look at a syntax for creating class objects that is new. If you create an object, you actually call a function that has the same name as the class: the *constructor*. By default there is a constructor which has no arguments, and does nothing. A constructor can for instance be used to initialize data members:

```
class Vector {  
private:  
    double x, y;  
public:
```

```

Vector( double userx,double usery ) {
    x = userx; y = usery;
}
/* ... */
} ;
/* ... */
Vector p1(1.,2.);

```

Other syntax for initialization:

```

class Vector {
private:
    double x,y;
public:
    Vector( double userx,double usery ) : x(userx),y(usery) {
    }
    /* ... */
} ;

```

If the data members follow a `public` directive, code outside the class can access the data members, both for getting and setting their values. This may be convenient for coding, but it's not a clean coding style. It's better to make data members `private`, and use `accessor` functions to get and set values.

```

class Vector {
private:
    double vx,vy;
public:
    Vector( double x,double y ) {
        vx = x; vy = y;
    };
    double x() { return vx; }; // 'accessor'
    double y() { return vy; };
} ;

int main() {
    Vector p1(1.,2.);
    cout << "p1 = " << p1.x() << "," << p1.y() << endl;
}

```

With the accessors, you have just seen a first example of a class `method`: a function that is only defined for objects of that class, and that have access to the private data of that object. Let's now look at more meaningful methods. For instance, for the `Vector` class you can define functions such as `length` and `angle`.

```

class Vector {
private:
    double vx,vy;
public:
    Vector( double x,double y ) {

```

```

    vx = x; vy = y;
};

double length() { return sqrt(vx*vx + vy*vy); };
double angle() { return 0.; /* something trig */ };
};

int main() {
    Vector p1(1.,2.);
    cout << "p1 has length " << p1.length() << endl;
}

```

By making these functions public, and the data members private,, you define an Application Programmer Interface (API) for the class:

- You are defining operations for that class; they are the only way to access the data of the object.
- The methods can use the data of the object, or alter it.
- The actual data of the object is not accessible from outside the object.

So far you have seen methods that use the data members of an object to return some quantity. It is also possible to alter the members. For instance, you may want to scale a vector by some amount:

```

class Vector {
/* ...
void scaleby( double a ) {
    vx *= a; vy *= a; }
/* ...
};

/* ...
Vector p1(1.,2.);
cout << "p1 has length " << p1.length() << endl;
p1.scaleby(2.);
cout << "p1 has length " << p1.length() << endl;

```

The methods you have seen so far only returned elementary datatypes. It is also possible to return an object, even from the same class. For instance, instead of scaling the members of a vector object, you could create a new object based on the scaled members:

```

class Vector {
/* ...
Vector scale( double a ) {
    return Vector( vx*a, vy*a ); }
/* ...
};

/* ...
cout << "p1 has length " << p1.length() << endl;
Vector p2 = p1.scale(2.);
cout << "p2 has length " << p2.length() << endl;

```

### 10.1.1 Default constructor

One of the more powerful ideas in C++ is that there can be more than one constructor. You will often be faced with this whether you want or not. The following code looks plausible:

```
Vector p1(1.,2.), p2;
cout << "p1 has length " << p1.length() << endl;
p2 = p1.scale(2.);
cout << "p2 has length " << p2.length() << endl;
```

but it will give an error message during compilation. The reason is that

```
Vector p;
```

calls the default constructor. Now that you have defined your own constructor, the default constructor no longer exists. So you need to define it explicitly:

```
Vector() {};
Vector( double x, double y ) {
    vx = x; vy = y;
};
```

### 10.1.2 Accessors

It is a good idea to make the data in an object private, to control outside access to it.

- Sometimes this private data is auxiliary, and there is no reason for outside access.
- Sometimes you do want outside access, but you want to precisely control by whom.

Accessor functions:

```
class thing {
private:
    float x;
public:
    float get_x() { return x; }
    void set_x(float v) { x = v; }
};
```

This has advantages:

- You can print out any time you get/set the value; great for debugging
- You can catch specific values: if  $x$  is always supposed to be positive, print an error (throw an exception) if nonpositive.

Better accessor:

```
class thing {
private:
    float x;
public:
    float &the_x() { return x; };
```

```
};

int main () {
    thing t;
    t.the_x() = 5;
    cout << t.the_x();
}
```

The function `the_x` returns a reference to the internal variable `x`.

If the internal variable is something indexable:

```
class thing {
private:
    vector<float> x;
public:
    operator[](int i) { return x[i]; }
};
```

You define the subscript operator `[]` for the object, in terms of indexing of the private vector.

### 10.1.3 Methods

Methods can be

- private, because they are only used internally;
- public, because they should be usable from outside a class object, for instance in the main program;
- protected, because they should be usable in derived classes (see section 10.3.1).

You can have multiple methods with the same name, as long as they can be distinguished by their argument types. This is known as *overloading*.

## 10.2 Inclusion relations between classes

The data members of an object can be of elementary datatypes, or they can be objects. For instance, if you write software to manage courses, each `Course` object will likely have a `Person` object, corresponding to the teacher.

```
class Course {
private:
    Person the_instructor;
    int year;
}
class Person {
    string name;
    ...
}
```

Designing objects with relations between them is a great mechanism for writing structured code, as it makes the objects in code behave like objects in the real world. The relation where one object contains another, is called a *has-a relation* between classes.

At this time, do exercises in section 39.2.

Sometimes a class can behave as if it includes an object of another class, while not actually doing so. For instance, a line segment can be defined from two points

```
class Segment {
private:
    Point starting_point, ending_point;
}
...
Segment somesegment;
Point somepoint = somesegment.get_the_end_point();
```

or from one point, and a distance and angle:

```
class Segment {
private:
    Point starting_point;
    float length,angle;
}
```

In both cases the code using the object is written as if the segment object contains two points. This illustrates how object-oriented programming can decouple the API of classes from their actual implementation.

Related to this decoupling, a class can also have two very different constructors.

```
class Segment {
private:
    // up to you how to implement!
public:
    Segment( Point start,float length,float angle )
    { .... }
    Segment( Point start,Point end ) { ... }
```

Depending on how you actually implement the class, the one constructor will simply store the defining data, and the other will do some conversion from the given data to the actually stored data.

This is another strength of object-oriented programming: you can change your mind about the implementation of a class without having to change the program that uses the class.

At this time, do the exercises in section 39.2

### 10.3 Inheritance

In addition to the has-a relation, there is the *is-a relation*, also called *inheritance*. Here one class is a special case of another class. Typically the object of the *derived class* (the special case) then also inherits

the data and methods of the *base class* (the general case).

```
class General {
protected: // note!
    int g;
public:
    void general_method() {};
};

class Special : public General {
public:
    void special_method() { g = ... };
};
```

How do you define a derived class?

- You need to indicate what the base class is:

```
class Special : public General { .... }
```

- The base class needs to declare its data members as `protected`: this is similar to `private`, except that they are visible to derived classes
- The methods of the derived class can then refer to data of the base class;
- Any method defined for the base class is available as a method for a derived class object.

The derived class has its own constructor, with the same name as the class name, but when it is invoked, it also calls the constructor of the base class. This can be the default constructor, but often you want to call the base constructor explicitly, with parameters that are describing how the special case relates to the base case. Example:

```
class General {
public:
    General( double x,double y ) {};
};

class Special : public General {
public:
    Special( double x ) : General(x,x+1) {};
};
```

- Multiple inheritance: an X is-a A, but also is-a B.  
This mechanism is somewhat dangerous.
- Virtual base class: you don't actually define a function in the base class, you only say 'any derived class has to define this function'.

### 10.3.1 Methods of base and derived classes

- A derived class can inherit a method from the base class.
- A derived class can define a method that the base class does not have.
- A derived class *override* a base class method:

```
class Base {
public:
```

```
    virtual f() { ... };  
};  
class Deriv : public Base {  
public:  
    virtual f() override { ... };  
};
```

Special syntax for *abstract method*:

```
class Base {  
public:  
    virtual void f() = 0;  
};  
class Deriv {  
public:  
    virtual void f() { ... };  
};
```

- The base class declares that every derived class has to define `f`
- The base class itself does not define this method: *abstract class*
- You can not make objects of the base class.

# Chapter 11

## Arrays

### 11.1 Traditional arrays

Static allocation, initialization.

```
// basic/array.cxx
{
    int numbers[] = {5, 4, 3, 2, 1};
    cout << numbers[3] << endl;
}
{
    int numbers[5]{5, 4, 3, 2, 1};
    cout << numbers[3] << endl;
}
{
    int numbers[5] = {2};
    cout << numbers[3] << endl;
}
```

Disadvantage: no test on bounds.

Exercise 11.1. Check whether an array is sorted.

Exercise 11.2. Find the maximum element in an array.

Also report at what index the maximum occurs.

Arrays can be passed to a subprogram, but the bound is unknown there.

```
// basic/array.cxx
void print_first_index( int ar[] ) {
    cout << "First index: " << ar[0] << endl;
}
{
    int numbers[] = {1, 4, 2, 5, 6};
    print_first_index(numbers);
}
```

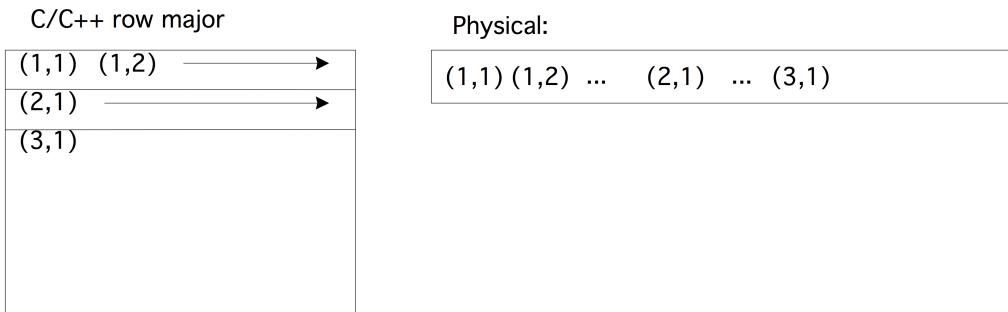
Exercise 11.3. Rewrite the above exercises where the sorting tester or the maximum finder is in a subprogram.

Subprograms can alter array elements. This was not possible with scalar arguments.

## 11.2 Multi-dimensional arrays

Declaration, pass to subprogram.

```
// array/contig.cxx
void print12( int ar[][6] ) {
    cout << "Array[1][2]: " << ar[1][2] << endl;
    return;
}
int array[5][6];
array[1][2] = 3;
print12(array);
```



Memory layout: multi-dimensional arrays are actually contiguous and linear.

```
// array/contig.cxx
void print06( int ar[][6] ) {
    cout << "Array[0][6]: " << ar[0][6] << endl;
    return;
}
int array[5][6];
array[1][0] = 35;
print06(array);
```

## 11.3 Other allocation mechanisms

A declaration

```
float ar[500];
```

is local to the scope it is in. This has some problems:

- Allocated on the stack; may lead to stack overflow.
- Can not be used as a class member:

```
class thing {
private:
    double array[ ???? ];
public:
    thing(int n) {
        array[ n ] ???? this does not work
    }
}
```

- Can not be returned from subprogram:

```
void make_array( double array[], int n ) {
    double array[n] ???? does not work
}
int main() {
    ...
    make_array(array, 100);
}
```

Use of `new` uses the equivalence of array and reference.

```
// array/arraynew.cxx
void make_array( int **new_array, int length ) {
    *new_array = new int[length];
}
int *the_array;
make_array(&the_array, 10000);
```

Since this is not scoped, you have to free the memory yourself:

```
// array/arraynew.cxx
class with_array{
private:
    int *array;
    int array_length;
public:
    with_array(int size) {
        array_length = size;
        array = new int[size];
    };
    ~with_array() {
        delete array;
    };
}
with_array thing_with_array(12000);
```

Notice how you have to remember the array length yourself? This is all much easier by using a `std::vector`. See <http://www.cplusplus.com/articles/37Mf92yv/>.

Deprecated use of `malloc`.

## 11.4 Vector class for arrays

Syntax:

```
#include <vector>
using namespace std;

vector<type> name(size)
```

where

- `vector` is a keyword,
- `type` (in angle brackets) is any elementary type or class name,
- `name` is up to you, and
- `size` is the (initial size of the array). This is an integer, or more precisely, a `size_t` parameter.

A vector behaves like an array:

```
vector<double> x(25);
x[1] = 3.14;
cout << x[2];
```

Initialization:

```
vector<int> odd_array{1,3,5,7,9};
vector<int> even_array = {0,2,4,6,8};
```

(This syntax requires compilation with the `-std=c++11` option.)

There is a second way of accessing elements:

```
vector<double> x(5);
x[5] = 1.; // will probably work
x.at(5) = 1.; // runtime error!
```

Safer, but also slower; see below.

Method `size` gives the size of the vector:

```
vector<char> words(37);
cout << words.size(); // will print 37
```

The vector class is a template class: the type that it uses (`int`, `float`) is not predetermined, but you can make a vector object out of whatever type you like.

### 11.4.1 Vector methods

- Get elements with `ar[3]` (zero-based indexing).
- Get elements, including bound checking, with `ar.at(3)`.
- Size: `ar.size()`.
- Other functions: `front`, `back`.

### 11.4.2 Vectors are dynamic

Use `push_back` to add elements at end.

```
vector<int> array(5);
array.push_back(35);
cout << array.size(); // is now 6 !
```

Other methods that change the size: `insert`, `erase`.

### 11.4.3 Vector assignment

The limitation that you couldn't create an array in an object still holds:

```
class witharray {
private:
    vector<int> the_array(????);
public:
    witharray( int n ) {
        thearray(???? n ?????);
    }
}
```

The following mechanism works:

```
class witharray {
private:
    vector<int> the_array;
public:
    witharray( int n ) {
        thearray = vector<int>(n);
    }
}
```

You could read this as

- `vector<int> the_array` declares a `int`-vector variable, and
- `thearray = vector<int>(n)` assigns an array to it.

However, technically, it actually does the following:

- The class object initially has a zero-size vector;
- the expression `vector<int>(n)` creates an anonymous vector of size `n`;
- which is then assigned to the variable `the_array`,
- so now you have an object with a vector of size `n` internally.

#### 11.4.4 Vectors and functions

##### 11.4.4.1 Vector as function return

You can have a vector as return type of a function:

```
vector<int> make_array(int n) {
    vector<int> x(n);
    x[0] = n;
    return x;
}
/* ... */
x1 = make_array(10);
```

##### 11.4.4.2 Pass vector to function

You can pass a vector to a function:

```
void print0( vector v ) {
    cout << v[0] << endl;
};
```

Vectors, like any argument, are passed by value, so the vector is actually copied into the function.

```
void set0( vector<float> &v, float x ) {
    v[0] = x;
}
/* ... */
vector<float> v(1);
v[0] = 3.5;
set0(v, 4.6);
cout << v[0] << endl;
```

This means you have to pass by reference:

**Exercise 11.4.** Write functions `random_vector` and `sort` to make the following main program work:

```
int length = 99;
vec<float> values = random_floats(length);
sort(values);
```

(This creates a vector of random values of a specified length, and then sorts it.)

#### 11.4.5 Dynamic size of vector

Writing

```
vector<int> iarray;
```

creates a vector of size zero. You can then

```
iarray.push_back(5);  
iarray.push_back(32);  
iarray.push_back(4);
```

to add elements to the vector, dynamically resizing it.

#### 11.4.6 Timing

Different ways of accessing a vector can have drastically different timing cost.

You can push elements into a vector:

```
vector<int> flex;  
point = std::chrono::system_clock::now();  
for (int i=0; i<LENGTH; i++)  
    flex.push_back(i);
```

If you allocate the vector statically, you can assign with at:

```
vector<int> stat(LENGTH);  
point = std::chrono::system_clock::now();  
for (int i=0; i<LENGTH; i++)  
    stat.at(i) = i;
```

or with subscript:

```
vector<int> stat(LENGTH);  
stat[0] = 0.;  
point = std::chrono::system_clock::now();  
for (int i=0; i<LENGTH; i++)  
    stat[i] = i;
```

You can also use new to allocate:

```
int *stat = new int[LENGTH];  
point = std::chrono::system_clock::now();  
for (int i=0; i<LENGTH; i++)  
    stat[i] = i;
```

Timings are partly predictable, partly surprising:

```
Flexible time: 2.445  
Static at time: 1.177  
Static assign time: 0.334  
Static assign time to new: 0.467
```

The increased time for new is a mystery.

## 11.5 Wrapping a vector in an object

You may want to create objects that contain a vector, for instance because you want to add some methods.

```
class printable {
private:
    vector<int> values;
public:
    printable(int n) {
        values = vector<int>(n);
    }
    string stringed() {
        string p("");
        for (int i=0; i<values.size(); i++)
            p += to_string(values[i])+" ";
        return p;
    }
};
```

Unfortunately this means you may have to recreate some methods:

```
int &at(int i) {
    return values.at(i);
};
```

## 11.6 Multi-dimensional cases

### 11.6.1 Matrix as vector of vectors

Multi-dimensional is harder with vectors:

```
vector<float> row(20);
vector<vector<float>> rows(10, row); // check on that >> syntax!
```

This is not contiguous.

### 11.6.2 Matrix class based on vector

Make sure you've studied classes; section ??.

You can make a ‘pretend’ matrix by storing a long enough `vector` in an object:

```
// array/matrixclass.cxx
class matrix {
private:
    std::vector<double> the_matrix;
    int m,n;
```

```

public:
    matrix(int m,int n) {
        this->m = m; this->n = n;
        the_matrix.reserve(m*n);
    };
    void set(int i,int j,double v) {
        the_matrix[ i*n +j ] = v;
    };
    double get(int i,int j) {
        return the_matrix[ i*n +j ];
    };
    /* ... */
};

```

The syntax for `set` and `get` can be improved.

**Exercise 11.5.** Write a method `element` of type `double&`, so that you can write

```
A.element(2,3) = 7.24;
```

## 11.7 Exercises

**Exercise 11.6.** Program *bubble sort*: go through the array comparing successive pairs of elements, and swapping them if the second is smaller than the first. After you have gone through the array, the largest element is in the last location. Go through the array again, swapping elements, which puts the second largest element in the one-before-last location. Et cetera.

**Exercise 11.7.** Pascal's triangle contains binomial coefficients:

Row	1:		1									
Row	2:		1	1								
Row	3:		1	2	1							
Row	4:		1	3	3	1						
Row	5:		1	4	6	4	1					
Row	6:		1	5	10	10	5	1				
Row	7:		1	6	15	20	15	6	1			
Row	8:		1	7	21	35	35	21	7	1		
Row	9:		1	8	28	56	70	56	28	8	1	
Row	10:		1	9	36	84	126	126	84	36	9	1

where

$$p_{rc} = \binom{r}{c} = \frac{r!}{c!(r-c)!}.$$

The coefficients can easily be computed from the recurrence

$$p_{rc} = \begin{cases} 1 & c \equiv 1 \vee c \equiv r \\ p_{r-1,c-1} + p_{r-1,c} & \text{otherwise} \end{cases}$$

- Write a class `pascal` so that `pascal(n)` is the object containing  $n$  rows of the above coefficients. Write a method `get(i, j)` that returns the  $(i, j)$  coefficient.
- Write a method `print` that prints the above display.
- Write a method `print(int m)` that prints a star if the coefficient modulo  $m$  is nonzero, and a space otherwise.

```
*
 *
 * *
 *
 *   *
 *
 *   *   *
 *
 *   *   *   *
 *
 *   *
 *
 *   *
```

- The object needs to have an array internally. The easiest solution is to make an array of size  $n \times n$ . When you have that code working, optimize your code to use precisely enough space for the coefficients.

**Exercise 11.8.** A knight on the chess board moves by going two steps horizontally or vertically, and one step either way in the orthogonal direction. Given a starting position, find a sequence of moves that brings a knight back to its starting position. Are there starting positions for which such a cycle doesn't exist?

**Exercise 11.9.** Put eight queens on a chessboard so that none threatens any other.

**Exercise 11.10.** From the ‘Keeping it REAL’ book, exercise 3.6 about Markov chains.

## Chapter 12

### Strings

#### 12.1 Basic string stuff

```
#include <string>
using namespace std;

// .. and now you can use 'string'
```

A *string* variable contains a string of characters.

```
string txt;
```

You can initialize the string variable (use `-std=c++11`), or assign it dynamically:

```
string txt{"this is text"};
string moretxt("this is also text");
txt = "and now it is another text";
```

Strings can be *concatenated*:

```
txt = txt1+txt2;
txt += txt3;
```

You can query the *size*:

```
int txtlen = txt.size();
```

or use subscripts:

```
cout << "The second character is <<" <<
    txt[1] << ">>" << endl;
```

Other methods for the vector class apply: `insert`, `empty`, `erase`, `push_back`, et cetera.

[http://en.cppreference.com/w/cpp/string/basic\\_string](http://en.cppreference.com/w/cpp/string/basic_string)

Exercise 12.1. Write a function to print out the digits of a number: 156 should print one five six. Use a vector or array of strings, containing the names of the digits.

Hint: it is easiest to generate the digits last-to-first. Then figure out how to print them reversed.

**Exercise 12.2.** Write a function to convert an integer to a string: the input 205 should give two hundred fifteen, et cetera.

**Exercise 12.3.** Write a pattern matcher, where a period . matches any one character, and x\* matches any number of ‘x’ characters.

For example:

- The string abc matches a.c but abbc doesn’t.
- The string abbc matches ab\*c, as does ac, but abzbc doesn’t.

## 12.2 Conversion

`to_string`

# Chapter 13

## Input/output

### 13.1 Formatted output

Normally, output of numbers takes up precisely the space that it needs:

```
cout << "Unformatted:" << endl;
for (int i=1; i<200000000; i*=10)
    cout << "Number: " << i << endl;
cout << endl;
```

Unformatted:  
Number: 1  
Number: 10  
Number: 100  
Number: 1000  
Number: 10000  
Number: 100000  
Number: 1000000  
Number: 10000000  
Number: 100000000

You can specify the number of positions, and the output is right aligned in that space by default:

```
cout << "Width is 6:" << endl;
for (int i=1; i<200000000; i*=10)
    cout << "Number: " << setw(6) << i << endl;
cout << endl;
```

(Only applies to immediately following number)

Width is 6:  
Number: 1  
Number: 10  
Number: 100  
Number: 1000  
Number: 10000  
Number: 100000

## 13. Input/output

---

```
Number: 1000000
Number: 10000000
Number: 100000000
```

Normally, padding is done with spaces, but you can specify other characters:

```
cout << "Padding:" << endl;
for (int i=1; i<200000000; i*=10)
    cout << "Number: "
        << left << setfill('.)' << setw(6) << i << endl;
```

Note: single quotes denote characters, double quotes denote strings.

Note: many of these output modifiers need

```
#include <iomanip>
```

```
Padding:
Number: .....1
Number: ....10
Number: ...100
Number: ..1000
Number: .10000
Number: 100000
Number: 1000000
Number: 10000000
Number: 100000000
```

Instead of right alignment you can do left:

```
cout << "Padding:" << endl;
for (int i=1; i<200000000; i*=10)
    cout << "Number: "
        << left << setfill('.)' << setw(6) << i << endl;
cout << endl;
```

```
Padding:
Number: 1.....
Number: 10....
Number: 100...
Number: 1000..
Number: 10000.
Number: 100000
Number: 1000000
Number: 10000000
Number: 100000000
```

Finally, you can print in different number bases than 10:

```
cout << "Base 16:" << endl;
cout << setbase(16) << setfill(' ');
for (int i=0; i<16; i++) {
    for (int j=0; j<16; j++)
        cout << i*16+j << " ";
    cout << endl;
}
cout << endl;

Base 16:
0 1 2 3 4 5 6 7 8 9 a b c d e f
10 11 12 13 14 15 16 17 18 19 1a 1b 1c 1d 1e 1f
20 21 22 23 24 25 26 27 28 29 2a 2b 2c 2d 2e 2f
30 31 32 33 34 35 36 37 38 39 3a 3b 3c 3d 3e 3f
etc
```

**Exercise 13.1.** Make the above output more nicely formatted:

```
00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f
10 11 12 13 14 15 16 17 18 19 1a 1b 1c 1d 1e 1f
20 21 22 23 24 25 26 27 28 29 2a 2b 2c 2d 2e 2f
```

**Exercise 13.2.** Use integer output to print fixed point numbers aligned on the decimal:

```
1.345
23.789
456.1234
```

Use four spaces for both the integer and fractional part.

## 13.2 Floating point output

Use `setprecision` to set the number of digits before and after decimal point:

```
x = 1.234567;
for (int i=0; i<10; i++) {
    cout << setprecision(4) << x << endl;
    x *= 10;
}

1.235
12.35
123.5
1235
1.235e+04
1.235e+05
1.235e+06
```

## 13. Input/output

---

```
1.235e+07  
1.235e+08  
1.235e+09
```

(Notice the rounding)

Fixed precision applies to fractional part:

```
cout << "Fixed precision applies to fractional part:" << endl;  
x = 1.234567;  
cout << fixed;  
for (int i=0; i<10; i++) {  
    cout << setprecision(4) << x << endl;  
    x *= 10;  
}  
  
1.2346  
12.3457  
123.4567  
1234.5670  
12345.6700  
123456.7000  
1234567.0000  
12345670.0000  
123456700.0000  
1234567000.0000
```

Combine width and precision:

```
x = 1.234567;  
cout << fixed;  
for (int i=0; i<10; i++) {  
    cout << setw(10) << setprecision(4) << x << endl;  
    x *= 10;  
}  
  
1.2346  
12.3457  
123.4567  
1234.5670  
12345.6700  
123456.7000  
1234567.0000  
12345670.0000  
123456700.0000  
1234567000.0000
```

```
cout << "Combine width and precision:" << endl;
x = 1.234567;
cout << scientific;
for (int i=0; i<10; i++) {
    cout << setw(10) << setprecision(4) << x << endl;
    x *= 10;
}

Combine width and precision:
1.2346e+00
1.2346e+01
1.2346e+02
1.2346e+03
1.2346e+04
1.2346e+05
1.2346e+06
1.2346e+07
1.2346e+08
1.2346e+09
```

### 13.3 Saving and restoring settings

```
ios::fmtflags old_settings = cout.flags();

cout.flags(old_settings);

int old_precision = cout.precision();

cout.precision(old_precision);
```

### 13.4 File output

Streams are general: work the same for console out and file out.

```
#include <fstream>
```

Use:

```
ofstream file_out;
file_out.open("fio_example.out");
/* ... */
file_out << number << endl;
file_out.close();
```

```
ofstream file_out;
file_out.open("fio_binary.out",ios::binary);
```

### 13.4.1 Output your own classes

You have used statements like:

```
cout << ``My value is: `` << myvalue << endl;
```

How does this work? The ‘double less’ is an operator with a left operand that is a stream, and a right operand for which output is defined; the result of this operator is again a stream. Recursively, this means you can chain any number of applications of << together.

If you want to output a class that you wrote yourself, you have to define how the << operator deals with your class.

```
class container {
    /* ... */
    int value() const {
        /* ... */
    };
    /* ... */
    std::ostream &operator<<(std::ostream &os, const container &i) {
        os << "Container: " << i.value();
        return os;
    };
    /* ... */
    container eye(5);
    cout << eye << endl;
```

## 13.5 Input streams

Test, mostly for file streams: `is_eof` `is_open`

## Chapter 14

### References and addresses

#### 14.1 Reference

This section contains further facts about parameter passing. Make sure you study section 7.1 first.

Passing a variable to a routine passes the value; in the routine, the variable is local.

```
// basic/arraypass.cxx
void change_scalar(int i) { i += 1; }
```

You can indicate that this is unintended:

```
// basic/arraypass.cxx
/* This does not compile:
   void change_const_scalar(const int i) { i += 1; }
*/
```

If you do want to make the change visible in the *calling environment*, use a reference:

```
// basic/arraypass.cxx
void change_scalar_by_reference(int &i) { i += 1; }
```

There is no change to the calling program. (Some people find this bad, since you can not see from the use of a function whether it passes *by reference* or *by value*.)

Arrays are always pass by reference:

```
// basic/arraypass.cxx
void change_array_location( int ar[], int i ) { ar[i] += 1; }
int numbers[5];
numbers[2] = 3.;
change_array_location(numbers,2);
```

The old-style way of doing things:

```
// basic/arraypass.cxx
void change_scalar_old_style(int *i) { *i += 1; }
number = 3;
change_scalar_old_style(&number);
```

## 14.2 Reference to class members

Here is the naive way of returning a class member:

```
class Object {  
private:  
    SomeType thing;  
public:  
    SomeType get_thing() {  
        return thing; };  
};
```

The problem here is that the return statement makes a copy of `thing`, which can be expensive. Instead, it is better to return the member by *reference*:

```
SomeType &get_thing() {  
    return thing; };
```

The problem with this solution is that the calling program can now alter the private member. To prevent that, use a *const reference*:

```
class has_int {  
private:  
    int mine{1};  
public:  
    const int& get() { return mine; };  
    int& set() { return mine; };  
    void inc() { mine++; };  
};  
/* ... */  
has_int an_int;  
an_int.inc(); an_int.inc(); an_int.inc();  
cout << "Contained int is now: " << an_int.get() << endl;  
/* Compiler error: an_int.get() = 5; */  
an_int.set() = 17;  
cout << "Contained int is now: " << an_int.get() << endl;
```

Output of this fragment:

```
Contained int is now: 4  
Contained int is now: 17
```

## 14.3 rvalue references

See the chapter about obscure stuff; section [24.2.3](#).

## **Chapter 15**

### **Polymorphism**

#### **15.1 The basic idea**

Sometimes you want to have the same function name for two slightly different purposes. C++ allows you to define the same function twice, as long as their parameters are different enough.

For instance, here is the same ‘sum’ function defined for both integers and reals:



## **Chapter 16**

### **Memory**

#### **16.1 Memory and scope**

If a variable goes *out of scope*, its memory is deallocated.

Deallocating objects is slightly more complicated: an *object destructor* is called. gives:

Before the nested scope  
calling the constructor  
Inside the nested scope  
calling the destructor  
After the nested scope



## Chapter 17

### Pointers

#### 17.1 What is a pointer

The term pointer is used to denote a reference to a quantity. The reason that people like to use C++ and C as high performance languages is that pointers are actually memory addresses. So you're programming ‘close to the bare metal’ and are in fargoing control over what your program does.

#### 17.2 Pointers and addresses, C style

You have learned about variables, and maybe you have a mental concept of variables as ‘named memory locations’. That is not too far of: while you are in the (dynamic) scope of a variable, it corresponds to a fixed memory location.

Exercise 17.1. When can a variable correspond to a varying location, depending on when you look at it?

There is a mechanism of finding the actual address of a variable: you prefix its name by an ampersand. This address is integer-valued, but its range is actually greater than of the `int` type.

If you have an

```
int i;
```

then `&i` is the address of `i`.

An address is a (long) integer, denoting a memory address. Usually it is rendered in *hexadecimal* notation:

```
int i;
printf("address of i: %ld\n", (long)(&i));
printf(" same in hex: %x\n", (long)(&i));
```

You could just print out the address of a variable, which is sometimes useful for debugging. If you want to store the address, you need to create a variable of the appropriate type. This is done by taking a type and affixing a star to it.

The type of ‘`&i`’ is `int*`, pronounced ‘int-star’, or more formally: ‘pointer-to-int’.

You can create variables of this type:

17. Pointers

```
int i;  
int* addr = &i;
```

Now if you have have a pointer that refers to an int:

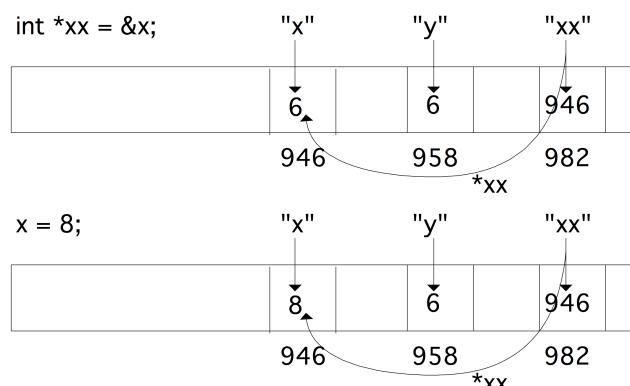
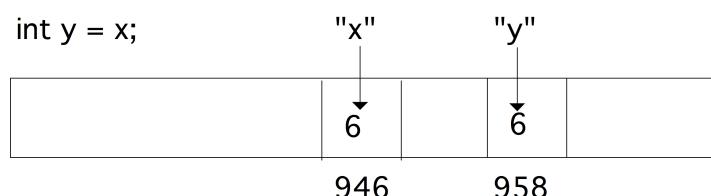
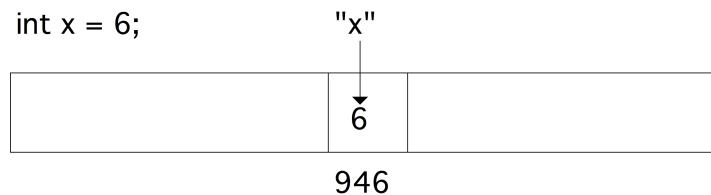
```
int i;  
int *iaddr = &i;
```

you can use (for instance `print`) that pointer, which gives you the address of the variable. If you want the value of the variable that the pointer points to, you need to *dereference* it.

Using `*addr` ‘dereferences’ the pointer: gives the thing it points to; the value of what is in the memory location.

```
int i;  
int* addr = &i;  
i = 5;  
cout << *addr;  
i = 6;  
cout << *addr;
```

This will print 5 and 6:



- `addr` is the address of `i`.
- You set `i` to 5; nothing changes about `addr`. This has the effect of writing 5 in the memory location of `i`.
- The first `cout` line dereferences `addr`, that is, looks up what is in that memory location.
- Next you change `i` to 6, that is, you write 6 in its memory location.
- The second `cout` looks in the same memory location as before, and now finds 6.

The syntax for declaring a pointer-to-sometype allows for a small variation, which indicates the two way you can interpret such a declaration.

Equivalent:

- `int* addr`: `addr` is an int-star, or
- `int *addr`: `*addr` is an int.

The notion `int* addr` is equivalent to `int *addr`, and semantically they are also the same: you could say that `addr` is an int-star, or you could say that `*addr` is an int.

### 17.2.1 Arrays and pointers

In section 11.1 you saw the treatment of static arrays in C++. Examples such as:

```
// basic/array.cxx
void print_first_index( int ar[] ) {
    cout << "First index: " << ar[0] << endl;
}
{
    int numbers[] = {1,4,2,5,6};
    print_first_index(numbers);
}
```

show that, even though array parameters are normally passed by value, that is through copying, array parameters can be altered. The reason for this is that there is no actual array type, and what is passed is a pointer to the first element of the array. So arrays are still passed by value, just not the ‘value of the array’, but the value of its location.

So you could pass an array like this:

```
// array/arraypass.cxx
void array_set_star( double *ar,int idx,double val) {
    ar[idx] = val;
}
array_set_star(array,2,4.2);
```

Array and memory locations are largely the same:

```
double array[5];
double *addr_of_second = &(array[1]);
array = (11,22,33,44,55);
cout << *addr_of_second;
```

`new` gives a something-star:

```
double *x;  
x = new double[27];
```

### 17.2.2 Pointer arithmetic

*pointer arithmetic* uses the size of the objects it points at:

```
double *addr_of_element = array;  
cout << *addr_of_element;  
addr_of_element = addr_of_element+1;  
cout << *addr_of_element;
```

Increment add size of the array element, 4 or 8 bytes, not one!

Exercise 17.2. Write a subroutine that sets the  $i$ -th element of an array, but using pointer arithmetic: the routine should not contain any square brackets.

### 17.2.3 Multi-dimensional arrays

After

```
double x[10][20];
```

a row  $x[3]$  is a `double*`, so is  $x$  a `double**`?

Was it created as:

```
double **x = new double*[10];  
for (int i=0; i<10; i++)  
    x[i] = new double[20];
```

No: multi-d arrays are contiguous.

### 17.2.4 Parameter passing

C++ style functions that alter their arguments:

```
void inc(int &i) { i += 1; }  
int main() {  
    int i=1;  
    inc(i);  
    cout << i << endl;  
    return 0;  
}
```

In C you can not pass-by-reference like this. Instead, you pass the address of the variable `i` by value:

```
void inc(int *i) { *i += 1; }
int main() {
    int i=1;
    inc(&i);
    cout << i << endl;
    return 0;
}
```

Now the function gets an argument that is a memory address: `i` is an int-star. It then increases `*i`, which is an int variable, by one.

**Exercise 17.3.** Write another version of the `swap` function:

```
void swap( /* something with i and j */ {
    /* your code */
}
int main() {
    int i=1, j=2;
    swap( /* something with i and j */ );
    cout << "check that i is 2: " << i << endl;
    cout << "check that j is 1: " << j << endl;
    return 0;
}
```

### 17.2.5 Allocation

In section 11.1 you learned how to create arrays that are local to a scope:

```
if ( something ) {
    double ar[25];
} else {
    double ar[26];
}
ar[0] = // there is no array!
```

The array `ar` is created depending on if the condition is true, but after the conditional it disappears again. The mechanism of using `new` (section 11.3) allows you to allocate storage that transcends its scope:

```
double *array;
if (something) {
    array = new double[25];
} else {
    array = new double[26];
}
```

Memory allocated with `new` does not disappear when you leave a scope. Therefore you have to delete the memory explicitly:

```
delete(array);
```

#### 17.2.5.1 Malloc

The keywords `new` and `delete` are in the spirit of C style programming, but don't exist in C. Instead, you use `malloc`, which creates a memory area with a size expressed in bytes. Use the function `sizeof` to translate from types to bytes:

```
int n;
double *array;
array = malloc( n*sizeof(double) );
if (!array)
    // allocation failed!
```

#### 17.2.5.2 Allocation in a function

The mechanism of creating memory, and assigning it to a 'star' variable can be used to allocate data in a function and return it from the function.

```
void make_array( double **a, int n ) {
    *a = new double[n];
}
int main() {
    double *array;
    make_array(&array,17);
}
```

Note that this requires a 'double-star' or 'star-star' argument:

- The variable `a` will contain an array, so it needs to be of type `double*`;
- but it needs to be passed by reference to the function, making the argument type `double**`;
- inside the function you then assign the new storage to the `double*` variable, which is `*a`.

Tricky, I know.

#### 17.2.6 Memory leaks

Pointers can lead to a problem called *memory leaking*: there is memory that you have reserved, but you have lost the ability to access it.

In this example:

```
double *array = new double[100];
// ...
array = new double[105];
```

memory is allocated twice. The memory that was allocated first is never released, because in the intervening code another pointer to it may have been set. However, if that doesn't happen, the memory is both allocated, and unreachable. That's what memory leaks are about.

### 17.3 Safer pointers in C++

Section 17.2.6 showed how memory can become unreachable. The *C++11* standard has mechanisms that can help solve this problem<sup>1</sup>.

A ‘shared pointer’ is a pointer that keeps count of how many times the object is pointed to. If one of these pointers starts pointing elsewhere, the shared pointer decreases this ‘reference count’. If the reference count reaches zero, the object is deallocated or destroyed.

```
#include <memory>

auto array = std::shared_ptr<double>( new double[100] );
```

As an illustration:

```
cout << "set pointer1" << endl;
auto thing_ptr1 = shared_ptr<thing>( new thing );
cout << "overwrite pointer" << endl;
thing_ptr1 = nullptr;
```

creates an object and overwrites the pointer, causing the object to be destroyed:

```
set pointer1
calling constructor
overwrite pointer
calling destructor
```

Illustrating how the object is not destroyed until all references are gone:

```
cout << "set pointer2" << endl;
auto thing_ptr2 = shared_ptr<thing>( new thing );
cout << "set pointer3 by copy" << endl;
auto thing_ptr3 = thing_ptr2;
cout << "overwrite pointer2" << endl;
thing_ptr2 = nullptr;
cout << "overwrite pointer3" << endl;
thing_ptr3 = nullptr;
```

gives:

```
set pointer2
calling constructor
set pointer3 by copy
overwrite pointer2
overwrite pointer3
calling destructor
```

---

1. A mechanism along these lines already existed in the ‘weak pointer’, but it was all but unusable.



# Chapter 18

## Prototypes

### 18.1 Prototypes for functions

In most of the programs you have written in this course, you put any functions or classes above the main program, so that the compiler could inspect the definition before it encountered the use. However, the compiler does not actually need the whole definition, say of a function: it is enough to know its name, the types of the input parameters, and the return type.

Such a minimal specification of a function is known as function *prototype*; for instance

```
int tester(float);
```

A first use of prototypes is *forward declaration*:

```
int f(int);
int g(int i) { return f(i); }
int f(int i) { return g(i); }
```

Prototypes are useful if you spread your program over multiple files. You would put your functions in one file:

```
// file: def.cxx
int tester(float x) {
    ....
}
```

and the main program in another:

```
// file : main.cxx
int tester(float);

int main() {
    int t = tester(...);
    return 0;
}
```

Or you could use your function in multiple programs and you would have to write it only once.

### 18.1.1 Header files

Even better than writing the prototype every time you need the function is to have a *header file*:

```
// file: def.h
int tester(float);
```

The definitions file would include this:

```
// file: def.cxx
#include "def.h"
int tester(float x) {
    ....
}
```

and so does the main program

```
// file : main.cxx
#include "def.h"

int main() {
    int t = tester(...);
    return 0;
}
```

Having a header file is an important safety measure:

- Suppose you change your function definition, changing its return type;
- The compiler will complain when you compile the definitions file;
- So you change the prototype in the header file;
- Now the compiler will complain about the main program, so you edit that too.

**Remark 1** *By the way, why does that compiler even recompile the main program, even though it was not changed? Well, that's because you used a makefile. See the tutorial.*

**Remark 2** *Header files were able to catch more errors in C than they do in C++. With polymorphism of functions, it is no longer an error to have*

```
// header.h
int somefunction(int);
```

and

```
#include "header.h"

int somefunction( float x ) { .... }
```

### 18.1.2 C and C++ headers

You have seen the following syntaxes for including header files:

```
#include <header.h>
#include "header.h"
```

The first is typically used for system files, with the second typically for files in your own project. There are some header files that come from the C standard library such as `math.h`; the idiomatic way of including them in C++ is

```
#include <cmath>
```

## 18.2 Global variables

If you have a variable that you want known everywhere, you can make it global:

```
int processnumber;
int main() {
    processnumber = // some system call
};
```

It is then defined in functions defined in your program file.

If your program has multiple files, you should not put ‘`int processnumber`’ in the other files, because that would create a new variable, that is only known to the functions in that file. Instead use:

```
extern int processnumber;
```

which says that the global variable `processnumber` is defined in some other file.

What happens if you put that variable in a *header file*? Since the *preprocessor* acts as if the header is textually inserted, this again leads to a separate global variable per file. The solution then is more complicated:

```
//file: header.h
#ifndef HEADER_H
#define HEADER_H
#ifndef EXTERN
#define EXTERN extern
#endif
EXTERN int processnumber
#endif

//file: aux.cc
#include "header.h"

//file: main.cc
#define EXTERN
```

```
#include "header.h"
```

This also prevents recursive inclusion of header files.

### 18.3 Prototypes for class methods

Header file:

```
class something {
public:
    double somedo(vector);
};
```

Implementation file:

```
double something::somedo(vector v) {
    .... something with v ....
};
```

### 18.4 Header files and templates

The use of *templates* often make separate compilation impossible: in order to compile the templated definitions the compiler needs to know with what types they will be used.

### 18.5 Namespaces and header files

Never put using namespace in a header file.

# Chapter 19

## Namespaces

### 19.1 Solving name conflicts

In section 11.4 you saw that the C++ library comes with a `vector` class, that implements dynamic arrays. You say

```
std::vector<int> bunch_of_ints;
```

and you have an object that can store a bunch of ints. And if you use such vectors often, you can save yourself some typing by having

```
using namespace std;
```

somewhere high up in your file, and write

```
vector<int> bunch_of_ints;
```

in the rest of the file.

But what if you are writing a geometry package, which includes a `vector` class? Is there confusion with the **STL!** (**STL!**) `vector` class? There would be if it weren't for the phenomenon *namespace*, which acts as a disambiguating prefix for classes, functions, variables.

You have already seen namespaces in action when you wrote `std::vector`: the ‘`std`’ is the name of the namespace. You can make your own namespace by writing

```
namespace a_namespace {  
    // definitions  
    class an_object {  
    };  
}
```

so that you can write

```
a_namespace::an_object myobject();
```

or

```
using namespace a_namespace;
an_object myobject();
```

or

```
using a_namespace::an_object;
an_object myobject();
```

### 19.1.1 Namespace header files

If your namespace is going to be used in more than one program, you want to have it in a separate source file, with an accompanying header file:

```
#include "geolib.h"
using namespace geometry;
```

The header would contain the normal function and class headers, but now inside a named namespace:

```
namespace geometry {
    class point {
        private:
            double xcoord,ycoord;
        public:
            point() {};
            point( double x,double y );
            double x();
            double y();
        };
        class vector {
        private:
            point from,to;
        public:
            vector( point from,point to);
            double size();
        };
}
```

and the implementation file would have the implementations, in a namespace of the same name:

```
namespace geometry {
    point::point( double x,double y ) {
        xcoord = x; ycoord = y; }
    double point::x() { return xcoord; } // 'accessor'
    double point::y() { return ycoord; }
    vector::vector( point from,point to) {
        this->from = from; this->to = to;
    };
}
```

```
double vector::size() {
    double
        dx = to.x()-from.x(), dy = to.y()-from.y();
    return sqrt( dx*dx + dy*dy );
};
```



## Chapter 20

### Preprocessor

In your source files you have seen lines starting with a hash sign, like

```
#include <iostream>
```

Such lines are interpreted by the *C preprocessor*.

Your source file is transformed to another source file, in a source-to-source translation stage, and only that second file is actually compiled by the *compiler*. In the case of an `#include` statement, the preprocessing stage takes form of literally inserting another file, here a *header file* into your source.

There are more sophisticated uses of the preprocessor.

#### 20.1 Textual substitution

Suppose your program has a number of arrays and loop bounds that are all identical. To make sure the same number is used, you can create a variable, and pass that to routines as necessary.

```
void dosomething(int n) {
    for (int i=0; i<n; i++) ....
}

int main() {
    int n=100000;

    double array[n];

    dosomething(n);
```

You can also use a *preprocessor macro*:

```
#define N 100000
void dosomething() {
    for (int i=0; i<N; i++) ....
}

int main() {
```

```
double array[N];
```

```
dosomething();
```

It is traditional to use all uppercase for such macros.

## 20.2 Parametrized macros

Instead of simple text substitution, you can have *parametrized preprocessor macros*

```
#define CHECK_FOR_ERROR(i) if (i!=0) return i  
...  
ierr = some_function(a,b,c); CHECK_FOR_ERROR(ierr);
```

When you introduce parameters, it's a good idea to use lots of parentheses:

```
// the next definition is bad!  
#define MULTIPLY(a,b) a*b  
...  
x = MULTIPLY(1+2, 3+4);
```

Better

```
#define MULTIPLY(a,b) (a)*(b)  
...  
x = MULTIPLY(1+2, 3+4);
```

Another popular use of macros is to simulate multi-dimensional indexing:

```
#define INDEX2D(i,j,n) (i)*(n)+j  
...  
double array[m,n];  
for (int i=0; i<m; i++)  
    for (int j=0; j<n; j++)  
        array[ INDEX2D(i,j,n) ] = ...
```

Exercise 20.1. Write a macro that simulates 1-based indexing:

```
#define INDEX2D1BASED(i,j,n) ????  
...  
double array[m,n];  
for (int i=1; i<=m; i++)  
    for (int j=n; j<=n; j++)  
        array[ INDEX2D1BASED(i,j,n) ] = ...
```

## 20.3 Conditionals

There are a couple of *preprocessor conditions*.

### 20.3.1 Check on a value

The `#if` macro tests on nonzero. A common application is to temporarily remove code from compilation:

```
#if 0
    bunch of code that needs to
        be disabled
#endif
```

### 20.3.2 Check for macros

The `#ifdef` test tests for a macro being defined. Conversely, `#ifndef` tests for a macro not being defined. For instance,

```
#ifndef N
#define N 100
#endif
```

Why would a macro already be defined? Well you can do that on the compile line:

```
icpc -c file.cc -DN=500
```

Another application for this test is in preventing recursive inclusion of header files; see section [18.2](#).



## Chapter 21

### Templates

Sometimes you want a function or a class based on more than one different datatypes. For instance, in chapter 11 you saw how you could create an array of ints as `vector<int>` and of doubles as `vector<double>`. Here you will learn the mechanism for that.

Basically, you want the type name to be a variable. Syntax:

```
template <typename yourtypevariable>
// ... stuff with yourtypevariable ...
```

Definition:

```
template<typename T>
void function(T var) { cout << var << endl; }
```

Usage:

```
int i; function(i);
double x; function(x);
```

and the code will behave as if you had defined `function` twice, once for `int` and once for `double`.

Exercise 21.1. Machine precision, or ‘machine epsilon’, is sometimes defined as the smallest number  $\epsilon$  so that  $1 + \epsilon > 1$  in computer arithmetic.

Write a templated function `epsilon` so that the following code prints out the values of the machine precision for the `float` and `double` type respectively:

```
float float_eps;
epsilon(float_eps);
cout << "For float, epsilon is " << float_eps << endl;

double double_eps;
epsilon(double_eps);
cout << "For double, epsilon is " << double_eps << endl;
```

## 21.1 Templating over non-types

THESE EXAMPLES ARE NOT GOOD.

See: <https://www.codeproject.com/Articles/257589/An-Idiots-Guide-to-Cplusplus-Templating>

Templating over integral types, not double.

The templated quantity is a value:

```
template<int s>
std::vector<int> svector(s);
/* ... */
svector(3) threevector;
cout << threevector.size();
```

Exercise 21.2. Write a class that contains an array. The length of the array should be templated.

## Chapter 22

### Error handling

#### 22.1 General discussion

Sources or errors:

- Array indexing. See section 11.4.
- Null pointers
- Division by zero and other numerical errors.

Guarding against errors.

- Check preconditions.
- Catch results.
- Check postconditions.

Error reporting:

- Message
- Total abort
- Exception

Assertions:

```
#include <cassert>
...
assert( bool )
```

assertions are omitted with optimization

Function return values

#### 22.2 Exception handling

*Throwing an exception* is one way of signalling an error or unexpected behaviour:

```
void do_something() {
    if ( oops )
        throw(5);
}
```

It now becomes possible to detect this unexpected behaviour by *catching* the exception:

```
throw {
    do_something();
} catch (int i) {
    cout << "doing something failed: error=" << i << endl;
}
```

You can throw integers to indicate an error code, a string with an actual error message. You could even make an error class: Sophisticated:

```
class MyError {
public :
    int error_no; string error_msg;
    MyError( int i, string msg )
        : error_no(i), error_msg(msg) {};
}

throw( MyError(27, "oops") ;

try {
    // something
} catch ( MyError &m ) {
    cout << "My error with code=" << m.error_no
    << " msg=" << m.error_msg << endl;
}
```

You can multiple `catch` statements to catch different types of errors:

```
try {
    // something
} catch ( int i ) {
    // handle int exception
} catch ( std::string c ) {
    // handle string exception
}
```

Catch all exceptions:

```
try {
    // something
} catch ( ... ) { // literally: three dots
    cout << "Something went wrong!" << endl;
}
```

- Functions can define what exceptions they throw:

```
void func() throw( MyError, std::string );
void funk() throw();
```

- Predefined exceptions: `bad_alloc`, `bad_exception`
- An exception handler can throw an exception; to rethrow the same exception use ‘`throw;`’ without arguments.
- Exceptions delete all stack data, but not `new` data.



## Chapter 23

### Standard Template Library

The C++ language has a *Standard Template Library* (STL), which contains functionality that is considered standard, but that is actually implemented in terms of already existing language mechanisms. The STL is enormous, so we just highlight a couple of parts.

You have already seen

- arrays (chapter 11),
- strings (chapter 12),
- streams (chapter 13).

#### 23.1 Containers

Vectors (section 11.4) and strings (chapter 12) are special cases of a STL *container*. Methods such as `push_back` and `insert` apply to all containers.

##### 23.1.1 Iterators

The container class has a subclass *iterator* that can be used to iterate through all elements of a container.

```
for (vector::iterator element=myvector.begin();
      element!=myvector.end(); elements++) {
    // do something with element
}
```

You would hope that, if `myvector` is a vector of `int`, `element` would be an `int`, but it is actually a pointer-to-`int`; section 17.2. So you could write

```
for (vector::iterator elt=myvector.begin();
      elt!=myvector.end(); elt++) {
    int element = *elt;
    // do something with element
}
```

This looks cumbersome, and you can at least simplify it by letting C++ deduce the type:

```
for (auto elt=myvector.begin(); ..... ) {  
    ....  
}
```

In the C++11/14 standard the iterator notation has been simplified to *range-based iteration*:

```
for ( int element : myvector ) {  
    ...  
}
```

## 23.2 Complex numbers

```
#include <complex>  
complex<float> f;  
f.re = 1.; f.im = 2.;  
  
complex<double> d(1., 3.);
```

Math operator like `+`, `*` are defined, as are math functions.

## 23.3 About the ‘using’ keyword

Only use this internally, not in header files that the user sees.

## Chapter 24

### Obscure stuff

#### 24.1 Casts

In C++, constants and variables have clear types. For cases where you want to force the type to be something else, there is the *cast* mechanism. With a cast you tell the compiler: treat this thing as such-and-such a type, no matter how it was defined.

##### 24.1.1 Casting constants

One use of casting is to convert constants to a ‘larger’ type. For instance, allocation does not use integers but `size_t`.

```
int hundredk = 100000;
int overflow;
overflow = hundredk*hundredk;
/* ... */
size_t bignum = static_cast<size_t>(hundredk)*hundredk;
```

C++ pointers are really memory addresses, with no type information to it. With a *cast* it becomes possible change your mind about what a pointer is.

The `static cast` is also useful for dealing with a *void pointer*.

The `static_cast` has the safety feature that the compiler will complain if the conversion is not possible.

##### 24.1.2 Dynamic cast

If we have a pointer to a derived object, stored in a pointer to a base class object, it’s possible to turn it safely into a derived pointer again:

```
derived_object *derived_pointer;
basic_class *basic_pointer;
derived_pointer = dynamic_cast<derived_object*>(basic_pointer);
if (derived_pointer==nullptr)
    // cast failed
```

Using a `static_cast` here would lead to a compiler error.

### 24.1.3 Legacy mechanism

The syntax ‘open parenthesis, type, closing parenthesis’ means:

- take whatever you have here,
- and interpret it as the specified type.

Example:

```
int i[2];
double *point_at_real = (double*)i;
cout << "Print two integers as double: " << *point_at_real << endl;
```

This is very dangerous. It is also impossible to search for such a thing in your editor. Please use the mechanisms above.

## 24.2 lvalue vs rvalue

The terms ‘lvalue’ and ‘rvalue’ sometimes appear in compiler error messages.

```
int foo() {return 2; }

int main()
{
    foo() = 2;

    return 0;
}

# gives:
test.c: In function 'main':
test.c:8:5: error: lvalue required as left operand of assignment
```

See the ‘lvalue’ and ‘left operand’? To first order of approximation you’re forgiven for thinking that an *lvalue* is something on the left side of an assignment. The name actually means ‘locator value’: something that’s associated with a specific location in memory. Thus an lvalue is, also loosely, something that can be modified.

An *rvalue* is then something that appears on the right side of an assignment, but is really defined as everything that’s not an lvalue. Typically, rvalues can not be modified.

The assignment `x=1` is legal because a variable `x` is at some specific location in memory, so it can be assigned to. On the other hand, `x+1=1` is not legal, since `x+1` is at best a temporary, therefore not at a specific memory location, and thus not an lvalue.

Less trivial examples:

```
int foo() { x = 1; return x; }
int main() {
    foo() = 2;
```

```
}
```

is not legal because `foo` does not return an lvalue. However,

```
class foo {
private:
    int x;
public:
    int &xfoo() { return x; }
};

int main() {
    foo x;
    x.xfoo() = 2;
```

is legal because the function `xfoo` returns a reference to the non-temporary variable `x` of the `foo` object.

Not every lvalue can be assigned to: in

```
const int a = 2;
```

the variable `a` is an lvalue, but can not appear on the left hand side of an assignment.

### 24.2.1 Conversion

Most lvalues can quickly be converted to rvalues:

```
int a = 1;
int b = a+1;
```

Here `a` first functions as lvalue, but becomes an rvalue in the second line.

The ampersand operator takes an lvalue and gives an rvalue:

```
int i;
int *a = &i;
&i = 5; // wrong
```

### 24.2.2 References

The ampersand operator yields a reference. It needs to be assigned from an lvalue, so

```
std::string &s = std::string(); // wrong
```

is illegal. The type of `s` is an ‘lvalue reference’ and it can not be assigned from an rvalue.

On the other hand

```
const std::string &s = std::string();
```

works, since `s` can not be modified any further.

### 24.2.3 Rvalue references

A new feature of C++ is intended to minimize the amount of data copying through *move semantics*.

Consider a copy assignment operator

```
BigThing& operator=( const BigThing &other ) {
    BigThing tmp(other); // standard copy
    std::swap( /* tmp data into my data */ );
    return *this;
};
```

This calls a copy constructor and a destructor on `tmp`. (The use of a temporary makes this safe under exceptions. The `swap` method never throws an exception, so there is no danger of half-copied memory.)

However, if you assign

```
thing = BigThing(stuff);
```

Now a constructor and destructor is called for the temporary rvalue object on the right-hand side.

Using a syntax that is new in C++, we create an *rvalue reference*:

```
BigThing& operator=( BigThing &&other ) {
    swap( /* other into me */ );
    return *this;
}
```

## **Chapter 25**

### **More exercises**

#### **25.1 Practice**

The website <http://www.codeforwin.in/2015/05/if-else-programming-practice.html> lists the following exercises for conditional:

1. Write a C program to find maximum between two numbers.
2. Write a C program to find maximum between three numbers.
3. Write a C program to check whether a number is even or odd.
4. Write a C program to check whether a year is leap year or not.
5. Write a C program to check whether a number is negative, positive or zero.
6. Write a C program to check whether a number is divisible by 5 and 11 or not.
7. Write a C program to count total number of notes in given amount.
8. Write a C program to check whether a character is alphabet or not.
9. Write a C program to input any alphabet and check whether it is vowel or consonant.
10. Write a C program to input any character and check whether it is alphabet, digit or special character.
11. Write a C program to check whether a character is uppercase or lowercase alphabet.
12. Write a C program to input week number and print week day.
13. Write a C program to input month number and print number of days in that month.
14. Write a C program to input angles of a triangle and check whether triangle is valid or not.
15. Write a C program to input all sides of a triangle and check whether triangle is valid or not.
16. Write a C program to check whether the triangle is equilateral, isosceles or scalene triangle.
17. Write a C program to find all roots of a quadratic equation.
18. Write a C program to calculate profit or loss.

#### **25.2 cplusplus**

<http://wwwcplusplus.com/forum/articles/12974/>

Dungeon crawl.

#### **25.3 world best learning center**

[http://www.worldbestlearningcenter.com/index\\_files/cpp-tutorial-variables\\_datatypes\\_exercises.htm](http://www.worldbestlearningcenter.com/index_files/cpp-tutorial-variables_datatypes_exercises.htm)



## **PART III**

### **FORTRAN**



## Chapter 26

### Basics of Fortran

Fortran is an old programming language, dating back to the 1950s, and the first ‘high level programming language’ that was widely used. In a way, the fields of programming language design and compiler writing started with Fortran, rather than this language being based on established fields. Thus, the design of Fortran has some idiosyncracies that later designed languages have not adopted. Many of these are now ‘deprecated’ or simply inadvisable. Fortunately, it is possible to write Fortran in a way that is every bit as modern and sophisticated as other current languages.

In this part of our book, we will teach you safe practices for writing Fortran. Occasionally we will not mention practices that you will come across in old Fortran codes, but that we would not advise you taking up. While our exposition of Fortran can stand on its own, we will in places point out explicitly differences with C++.

#### 26.1 Main program

Fortran does not use curly brackets to delineat blocks, instead you will find `end` statements. The very first one appears right when you start writing your program: a Fortran program needs to start with a `Program` line, and end with `End Program`. The program needs to have a name on both lines:

```
Program SomeProgram
    ! stuff goes here
End Program SomeProgram
```

and you can not use that name for any entities in the program.

##### 26.1.1 Program structure

Unlike C++, Fortran can not mix variable declarations and executable statements, so both the main program and any subprograms have roughly a structure:

```
Program foo
    < declarations >
    < statements >
End Program foo
```

(The *emacs* editor will supply the block type and name if you supply the ‘end’ and hit the TAB or RETURN key; see section 2.1.1.)

### 26.1.2 Statements

Let's say a word about layout. Fortran has a 'one line, one statement' principle.

- As long as a statement fits on one line, you don't have to terminate it explicitly with something like a semicolon:

```
x = 1  
y = 2
```

- If you want to put two statements on one line, you have to terminate the first one:

```
x = 1; y = 2
```

- If a statement spans more than one line, all but the first line need to have an explicit *continuation character*, the ampersand:

```
x = very &  
long &  
expression
```

(This is different between *free format* and *fixed format*, where it's the lines after the first that are marked a continuation, but we don't teach that here.)

### 26.1.3 Comments

Fortran knows only single-line *comments*, indicated by an exclamation point:

```
x = 1 ! set x to one
```

## 26.2 Variables

### 26.2.1 Data types

Fortran has a somewhat unusual treatment of data types: if you don't specify what data type a variable is, Fortran will deduce it from some default or user rules. This is a very dangerous practice, so we advocate putting a line

```
implicit none
```

immediately after any program or subprogram header.

Indicate number of bytes:

```
integer(2) :: i2  
integer(4) :: i4  
integer(8) :: i8  
  
real(4) :: r4  
real(8) :: r8  
real(16) :: r16
```

```
complex(8) :: c8
complex(16) :: c16
complex*32 :: c32
```

F08: `storage_size` reports number of bits.

F95: `bit_size` works on integers only.

`c_sizeof` reports number of bytes, requires `iso_c_binding` module.

### 26.2.2 Constants

Force a constant to be `real(8)`:

```
1.d0
```

### 26.2.3 Initialization

Variables can be initialized in their declaration:

```
integer :: i=2
real(4) :: x = 1.5
```

That this is done at compile time, leading to a common error:

```
subroutine foo()
    integer :: i=2
    cout << i << endl;
    i = 3
end subroutine foo
```

On the first subroutine call `i` is printed with its initialized value, but on the second call this initialization is not repeated, and the previous value of 3 is remembered.



## Chapter 27

### Loop constructs

#### 27.1 Loop types

Indexed do loop:

```
do i=1,10,2
    print *,i
end do
```

While:

```
do
    print *,i
    i = i*2
while (i<1000)
```

*F77 note:* Do not use label-terminated loops.

#### 27.2 Interruptions of the control flow

You can have a do loop without index or while test. In that case you need the exit statement to stop the iteration.

```
do
    x = randomvalue()
    if (x>.9) exit
    print *, "Nine out of ten exes agree"
end do
```

Skip rest of iteration:

```
do i=1,100
    if (isprime(i)) cycle
        ! do something with non-prime
    end do
```

Cycle and exit can apply to multiple levels, if the do-statements are labeled.

### 27.3 Implied do-loops

Iterate an item:

```
print *, (2*i, i=1, 20)
```

Multiple items:

```
print *, (2*i, 2*i+1, i=1, 20)
```

Nested:

```
print *, ( (i*j, i=1, 20), j=1, 20 )
```

Useful in I/O, especially of arrays.

# Chapter 28

## Scope, subprograms, modules

### 28.1 Scope

Fortran ‘has no curly brackets’: you not easily create nested scopes. The range between `do` and `end do` is not a scope.

However, see modules below [28.3](#).

### 28.2 Procedures

Programs can have procedures: parts of code that for some reason you want to separate from the main program. If you structure your code in a single file, this is the recommended structure:

```
Program foo
    < declarations>
    < executable statements >
    Contains
        < procedure definitions >
End Program foo
```

That is, procedures are placed after the main program statements, separated by a `contains` clause.

#### 28.2.1 Subroutines and functions

Fortran has two types of procedures:

- Subroutines, which are somewhat like `void` functions in C++: they can be used to structure the code, and they can only return information to the calling environment through their parameters.
- Functions, which are like C++ functions with a return type.

Both types have the same structure, which is roughly the same as of the main program:

```
subroutine foo( <parameters> )
    <variable declarations>
    <executable statements>
end subroutine foo
```

Exit from a procedure can happen two ways:

1. the flow of control reaches the end of the procedure body, or
2. execution is finished by an explicit `return` statement.

```
subroutine foo()
    print *, "foo"
    if (something) return
    print *, "bar"
end subroutine foo
```

The `return` statement is optional in the first case.

A subroutine is invoked with a `call` statement:

```
call foo()
```

### 28.2.2 Return results

While a subroutine can only return information through its parameters, a *function* procedure returns an explicit result:

```
logical function test(x)
    implicit none
    real :: x

    test = some_test_on(x)
    return ! optional, see above
end function test
```

You see that the result is not returned in the `return` statement, but rather through assignment to the function name. The `return` statement, as before, is optional and only indicates where the flow of control ends.

A function is not invoked with `call`, but rather through being used in an expression:

```
if (test(3.0) .and. something_else) ...
```

You now have the following cases to make the function known in the main program:

- If the function is in a `contains` section, its type is known in the main program.
- If the function is in a module (see section 28.3 below), it becomes known through a `use` statement.

*F77 note:* Without modules and `contains` sections, you need to declare the function type explicitly in the calling program.

### 28.2.3 Types of procedures

Procedures that are in the main program (or another type of program unit), separated by a `contains` clause, are known as *internal procedures*. This is as opposed to *module procedures*.

There are also *statement functions*, which are single-statement functions, usually to identify commonly used complicated expressions in a program unit. Presumably the compiler will *inline* them for efficiency.

The `entry` statement is so bizarre that I refuse to discuss it.

#### 28.2.4 Optional arguments

### 28.3 Modules

A module is a container for variables and functions, but it is not a class or object: there is only one instance.

```
Module FunctionsAndValues
    implicit none

    real(8),parameter :: pi = 3.14

    contains
        subroutine SayHi()
            print *, "Hi!"
        end subroutine SayHi

End Module FunctionsAndValues
```

Any routines come after the `contains`

A module is made available with the `use` keyword, which needs to go before the `implicit none`.

```
Program ModProgram
    use FunctionsAndValues
    implicit none

    print *, "Pi is:",pi
    call SayHi()

End Program ModProgram
```

If you compile a module, you will find a `.mod` file in your directory. (This is little like a `.h` file in C++.) If this file is not present, you can not use the module in another program unit.



## Chapter 29

### Structures, eh, types

The Fortran name for structures is `type` or *derived type*.

Type name / End Type block. Variable declarations inside the block

```
type mytype
    integer :: number
    character :: name
    real(4) :: value
end type mytype
```

Declare a typed object in the main program:

```
Type (mytype) :: typed_object,object2
```

Initialize with type name:

```
typed_object = mytype( 1, 'my_name', 3.7 )
object2 = typed_object
```

Access structure members with %

```
Type (mytype) :: typed_object
type_object%member = ....
```

```
type point
    real :: x,y
end type point

type(point) :: p1,p2
p1 = point(2.5, 3.7)

p2 = p1
print *,p2%x,p2%y
```



# Chapter 30

## Classes and objects

### 30.1 Classes

Fortran classes are based on `type` objects, a little like the analogy between C++ `struct` and `class` constructs.

New syntax for specifying methods.

You define a type as before, with its data members, but now the type has a `contains`.

```
% -*- latex -*-
%%%%%
%%% This TeX file is part of the course
%%% Introduction to Scientific Programming in C++/Fortran2003
%%% copyright 2017 Victor Eijkhout eijkhout@tacc.utexas.edu
%%%
%%% pointf.tex : arrays in Fortran
%%%
%%%%%
```

Pointers in C++ were largely the same as memory addresses (until you got to smart pointers). Fortran pointers on the other hand, are more abstract.

```
\Level 0 {Basic pointer operations}

\begin{block}{Pointers are aliases}
\begin{itemize}
\item Pointer points at an object
\item Access object through pointer
\item You can change what object the pointer points at.
\end{itemize}
\begin{verbatim}
real,pointer :: point_at_real
\end{verbatim}
\end{block}
```

Pointers could also be called ‘aliases’: they act like an alias for an object of elementary or derived data type. You can access the object through the alias. The difference with actually using the object, is that you can decide what object the pointer points at.

The `\index{pointer}` definition

```
\begin{verbatim}
real,pointer :: point_at_real
\end{verbatim}
```

defined a pointer that can point at a real variable.

```
\begin{block}{Setting the pointer}
  \begin{itemize}
    \item You have to declare that a variable is pointable:
  \begin{verbatim}
real,target :: x
\end{verbatim}
    \item Set the pointer with \verb+=>+ notation:
  \begin{verbatim}
point_at_real => x
\end{verbatim}
    \item Now using \n{point_at_real} is the same as using\n{x}.
  \end{itemize}
\end{block}
```

Pointers can not just point at anything: the thing pointed at needs to be declared as `\index{target}`

```
\begin{verbatim}
real,target :: x
\end{verbatim}
```

and you use the `\verb+=>+` operator to let a pointer point at a target:

```
\begin{verbatim}
point_at_real => x
\end{verbatim}
```

If you use a pointer, for instance to print it

```
\begin{verbatim}
print *,point_at_real
\end{verbatim}
```

it behaves as if you were using the value of what it points at.

```
\begin{block}{Pointer example}
  \verb+snippet{pointatreal}+
  %
\begin{enumerate}

```

```

\item The pointer points at~\n{x}, so the value of \n{x} is printed.
\item The pointer is set to point at~\n{y}, so its value is printed.
\item The value of \n{y} is changed, and since the pointer still
      points at~\n{y}, this changed value is printed.
\end{enumerate}
\end{block}

\begin{block}{Assign pointer from other pointer}
\begin{verbatim}
real,pointer :: point_at_real,also_point
point_at_real => x
also_point => point_at_real
\end{verbatim}
Now you have two pointers that point at~\n{x}.

\textbf{Very important to use the \n{=>}, otherwise strange
memory errors}
\end{block}

If you have two pointers
\begin{verbatim}
real,pointer :: point_at_real,also_point
\end{verbatim}
you can make the target of the one to also be the target of the other:
\begin{verbatim}
also_point => point_at_real
\end{verbatim}
This is not a pointer to a pointer: it assigns the target of the
right-hand side to be the target of the left-hand side.

\textbf{Using ordinary assignment does not work, and will give strange
memory errors.}

\begin{exercise}
Write a routine that accepts an array and a pointer, and on return
has that pointer pointing at the largest array element.
\end{exercise}
\begin{answer}
\verb+arraypointf+
\end{answer}

\Level 0 {Example: linked lists}

\begin{block}{Linked list}
\begin{itemize}
\item Linear data structure
\end{itemize}


```

```
\item more flexible for insertion~/ deletion
\item \ldots but slower in access
\end{itemize}
\end{block}
```

One of the standard examples of using pointers is the `\index{linked list}`. This is a dynamic one-dimensional structure that is more flexible than an array. Dynamically extending an array would require re-allocation, while in a list an element can be inserted.

```
\begin{exercise}
Using a linked list may be more flexible than using an array.
On the other hand, accessing an element in a linked list is
more expensive, both absolutely and as order-of-magnitude in the size
of the list.
```

Make this argument precise.  
\end{exercise}

```
\begin{block}{Linked list datatypes}
\begin{itemize}
\item Node: value field, and pointer to next node.
\item List: pointer to head node.
\end{itemize}
\verb|linklistf|
\end{block}
```

A list is based on a simple data structure, a node, which contains a value field and a pointer to another node.

By way of example, we create a dynamic list of integers, sorted by size. To maintain the sortedness, we need to append or insert nodes, as required.

Here are the basic definitions of a node, and a list which is basically a repository for the head node:

```
%  
\verb|linklistf|
```

```
\begin{block}{List initialization}
First element becomes the list head:  

\verb|listheadf|
\end{block}
```

Initially, the list is empty, meaning that the ‘head’ pointer is un-associated. The first time we add an element to the list, we create a node and assign it as the head of the list:

```
%  
\verbatimsnippet{listheadf}
```

```
\begin{block}{Attaching a node}
```

Keep the list sorted: new largest element attached at the end.

```
\verbatimsnippet{listattachf}  
\end{block}
```

Adding a value to a list can be done two ways. If the new element is larger than all elements in the list, a new node needs to be appended to the last one. Let’s assume we have managed to let `\n{current}` point at the last node of the list, then here is how to attaching a new node from it:

```
%  
\verbatimsnippet{listattachf}
```

```
\begin{block}{Inserting 1}
```

Find the insertion point:

```
\verbatimsnippet{listfindf}  
\end{block}
```

Inserting an element in the list is harder.

First of all, you need to find the two nodes,

`\n{previous}` and `\n{current}`, between which to insert the new node:

```
%
```

```
\verbatimsnippet{listfindf}
```

```
\begin{block}{Inserting 2}
```

The actual insertion requires rerouting some pointers:

```
%
```

```
\verbatimsnippet{listinsertf}
```

```
\end{block}
```

You define functions that accept the type as first argument, but instead of declaring the argument as `type`, you define it as `class`.

The members of the class object have to be accessed through the `%` operator.

```
subroutine set(p,xu,yu)  
  implicit none  
  class(point) :: p
```

```
real(8), intent(in) :: xu,yu
p%x = xu; p%y = yu
end subroutine set
```

Class objects are defined as `type` objects, just as if there were no class functions on them. The class functions are accessed as

```
object%function(arg1,arg2)
```

where the arguments do not include the `class` argument.

```
use PointClass
implicit none
type(Point) :: p1,p2

call p1%set(1.d0,1.d0)
call p2%set(4.d0,5.d0)
```

It is of course best to put the type definition and method definitions in a module, so that you can use it.

Mark methods as `private` so that they can only be used as part of the `type`:

```
Module PointClass
  private
contains
  subroutine setzero(p)
    implicit none
    class(point) :: p
    p%x = 0.d0 ; p%y = 0.d0
  end subroutine setzero
End Module PointClass
```

**Exercise 30.1.** Take the point example program and add a distance function.

# Chapter 31

## Arrays

### 31.1 Static arrays

The preferred way for specifying an array size is:

```
real(8), dimension(100) :: x,y
```

Such an array, with size explicitly indicated, is called a *static array* or *automatic array*. (See section 31.1.3 for dynamic arrays.)

Array indexing in Fortran is 1-based:

```
real :: x(8)
do i=1,8
... x(i) ...
```

Unlike C++, Fortran can specify the lower bound explicitly:

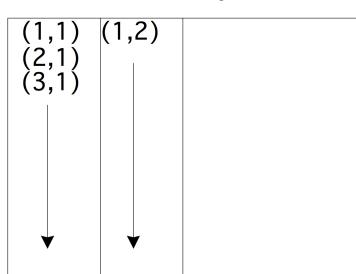
```
real :: x(-1:7)
do i=-1,7
... x(i) ...
```

Such array, as in C++, obey the scope: they disappear at the end of the program or subprogram.

#### 31.1.1 Multi-dimensional

```
integer,dimension(20,30) :: array
```

Fortran column major



Physical:

(1,1)	(2,1)	(3,1)	...	(1,2)	...
-------	-------	-------	-----	-------	-----

### 31.1.2 Query the size of an array

Various query functions:

- Total number of elements (in a specified dimension) can be queried with `size`:

```
integer :: x(8), y(5,4)
size(x)
size(y,2)
```

- Lower and upper bound:

```
Lbound(x)
Ubound(x)
```

### 31.1.3 Allocatable arrays

Static arrays are fine at small sizes. However, there are two main arguments against using them at large sizes.

- Since the size is explicitly stated, it makes your program inflexible, requiring recompilation to run it with a different problem size.
- Since they are allocated on the so-called *stack*, making them too large can lead to *stack overflow*.

A better strategy is to indicate the shape of the array, and use `allocate` to specify the size later, presumably in terms of run-time program parameters.

```
real(8), dimension(:), allocatable :: x,y
n = 100
allocate(x(n), y(n))
```

You can `deallocate` the array when you don't need the space anymore.

If you are in danger of running out of memory, it can be a good idea to add a `stat=ierror` clause to the `allocate` statement:

```
integer :: ierr
allocate( x(n), stat=ierr )
if ( ierr/=0 ) ! report error
```

Has an array been allocated:

```
Allocated( x ) ! returns logical
```

## 31.2 Arrays to subroutines

Subprogram needs to know the shape of an array, not the actual bounds:

```

real(8) function arraysum(x)
    implicit none
    real(8),intent(in),dimension(:) :: x
/* ...
do i=1,size(x)
    tmp = tmp+x(i)
end do
/* ...
Program ArrayComputations1D
use ArrayFunction
implicit none

real(8),dimension(:) :: x(N)
/* ...
print *, "Sum of one-based array:",arraysum(x)

```

### 31.3 Array slicing

### 31.4 Operating on an array

#### 31.4.1 Arithmetic operations

Between arrays of the same shape:

$$\begin{aligned} A &= B+C \\ D &= D \star E \end{aligned}$$

(where the multiplication is by element).

#### 31.4.2 Intrinsic functions

The following intrinsic functions are available for arrays:

`MaxVal` finds the maximum value in an array.

`MinVal` finds the minimum value in an array.

`Sum` returns the sum of all elements.

`Product` return the product of all elements.

`MaxLoc` returns the index of the maximum element.

`MinLoc` returns the index of the minimum element.

`MatMul` returns the matrix product of two matrices.

`DotProduct` returns the dot product of two arrays.

`Transpose` returns the transpose of a matrix.

`Cshift` rotates elements through an array.

**Exercise 31.1.** The 1-norm of a matrix is defined as the maximum sum of absolute values in any column:

$$\|A\|_1 = \max_j \sum_i |A_{ij}|$$

while the infinity-norm is defined as the maximum row sum:

$$\|A\|_{\infty} = \max_i \sum_j |A_{ij}|$$

Implement these functions using array intrinsics.

### 31.4.3 Restricting with where

where ( A<0 ) B = 0

Full form:

```
WHERE ( logical argument )
      sequence of array statements
ELSEWHERE
      sequence of array statements
END WHERE
```

## Chapter 32

### Pointers

Pointers in C++ were largely the same as memory addresses (until you got to smart pointers). Fortran pointers on the other hand, are more abstract.

#### 32.1 Basic pointer operations

- Pointer points at an object
- Access object through pointer
- You can change what object the pointer points at.

```
real,pointer :: point_at_real
```

Pointers could also be called ‘aliases’: they act like an alias for an object of elementary or derived data type. You can access the object through the alias. The difference with actually using the object, is that you can decide what object the pointer points at.

The `pointer` definition

```
real,pointer :: point_at_real
```

defined a pointer that can point at a real variable.

- You have to declare that a variable is pointable:

```
real,target :: x
```

- Set the pointer with `=>` notation:

```
point_at_real => x
```

- Now using `point_at_real` is the same as using `x`.

Pointers can not just point at anything: the thing pointed at needs to be declared as `target`

```
real,target :: x
```

and you use the `=>` operator to let a pointer point at a target:

```
point_at_real => x
```

If you use a pointer, for instance to print it

```
print *,point_at_real
```

it behaves as if you were using the value of what it points at.

```
real,target :: x,y  
real,pointer :: that_real  
  
x = 1.2  
y = 2.4  
that_real => x  
print *,that_real  
that_real => y  
print *,that_real  
y = x  
print *,that_real
```

1. The pointer points at `x`, so the value of `x` is printed.
2. The pointer is set to point at `y`, so its value is printed.
3. The value of `y` is changed, and since the pointer still points at `y`, this changed value is printed.

```
real,pointer :: point_at_real,also_point  
point_at_real => x  
also_point => point_at_real
```

Now you have two pointers that point at `x`.

### **Very important to use the =>, otherwise strange memory errors**

If you have two pointers

```
real,pointer :: point_at_real,also_point
```

you can make the target of the one to also be the target of the other:

```
also_point => point_at_real
```

This is not a pointer to a pointer: it assigns the target of the right-hand side to be the target of the left-hand side.

### **Using ordinary assignment does not work, and will give strange memory errors.**

Exercise 32.1. Write a routine that accepts an array and a pointer, and on return has that pointer pointing at the largest array element.

## 32.2 Example: linked lists

- Linear data structure
- more flexible for insertion / deletion

- ... but slower in access

One of the standard examples of using pointers is the *linked list*. This is a dynamic one-dimensional structure that is more flexible than an array. Dynamically extending an array would require re-allocation, while in a list an element can be inserted.

Exercise 32.2. Using a linked list may be more flexible than using an array. On the other hand, accessing an element in a linked list is more expensive, both absolutely and as order-of-magnitude in the size of the list.

Make this argument precise.

- Node: value field, and pointer to next node.
- List: pointer to head node.

```
type node
    integer :: value
    type(node),pointer :: next
end type node

type list
    type(node),pointer :: head
end type list

type(list) :: the_list
nullify(the_list%head)
```

A list is based on a simple data structure, a node, which contains a value field and a pointer to another node.

By way of example, we create a dynamic list of integers, sorted by size. To maintain the sortedness, we need to append or insert nodes, as required.

Here are the basic definitions of a node, and a list which is basically a repository for the head node:

```
type node
    integer :: value
    type(node),pointer :: next
end type node

type list
    type(node),pointer :: head
end type list

type(list) :: the_list
nullify(the_list%head)
```

First element becomes the list head:

```
allocate(new_node)
new_node%value = value
nullify(new_node%next)
```

```
the_list%head => new_node
```

Initially, the list is empty, meaning that the ‘head’ pointer is un-associated. The first time we add an element to the list, we create a node and assign it as the head of the list:

```
allocate(new_node)
new_node%value = value
nullify(new_node%next)
the_list%head => new_node
```

Keep the list sorted: new largest element attached at the end.

```
allocate(new_node)
new_node%value = value
nullify(new_node%next)
current%next => new_node
```

Adding a value to a list can be done two ways. If the new element is larger than all elements in the list, a new node needs to be appended to the last one. Let’s assume we have managed to let `current` point at the last node of the list, then here is how to attaching a new node from it:

```
allocate(new_node)
new_node%value = value
nullify(new_node%next)
current%next => new_node
```

Find the insertion point:

```
current => the_list%head ; nullify(previous)
do while ( current%value < value &
           .and. associated(current%next) )
    previous => current
    current => current%next
end do
```

Inserting an element in the list is harder. First of all, you need to find the two nodes, `previous` and `current`, between which to insert the new node:

```
current => the_list%head ; nullify(previous)
do while ( current%value < value &
           .and. associated(current%next) )
    previous => current
    current => current%next
end do
```

The actual insertion requires rerouting some pointers:

```
allocate(new_node)
new_node%value = value
```

```
new_node%next => current  
previous%next => new_node
```



# Chapter 33

## Input/output

### 33.1 Print to terminal

The simplest command for outputting data is `print`.

```
print *, "The result is", result
```

In its easiest form, you use the star to indicate that you don't care about formatting; after the comma you can then have any number of comma-separated strings and variables.

#### 33.1.1 Printing arrays

If you print a variable that is an array, all its elements will be printed, in *column-major* ordering if the array is multi-dimensional.

You can also control the printing of an array by using an *implicit do loop*:

```
print *, ( A(i,i), i=1, n)
```

#### 33.1.2 Formats

The default formatting uses quite a few positions for what can be small numbers. To indicate explicitly the formatting, for instance limiting the number of positions used for a number, or the whole and fractional part of a real number, you can use a format string.

```
print '(a6,3f5.3)', "Result", x, y, z
```

The format specifier is inside single quotes and parentheses, and consists of comma-separated specifications for a single item:

- `an` specifies a string of  $n$  characters. If the actual string is longer, it is truncated in the output.
- `in` specifies an integer of up to  $n$  digits. If the actual number takes more digits, it is rendered with asterisks.
- `fm.n` specifies a fixed point representation of a floating point number, with  $m$  total positions (including the decimal point) and  $n$  digits in the fractional part.
- `em.n` Exponent representation.

Putting a number in front of a single specifier indicates that it is to be repeated.

If you find yourself using the same format a number of times, you can give it a *label*:

```
print 10,"result:",x,y,z  
10 format(' (a6,3f5.3)' )
```

<https://www.obliquity.com/computer/fortran/format.html>

## 33.2 File I/O

Units. open close

## Chapter 34

### Array operations

#### 34.1 Loops without looping

In addition to ordinary do-loops, Fortran has mechanisms that save you typing, or can be more efficient in some circumstances.

- Slicing: if your loop assigns to an array from another array, you can use slice notation:

```
a(:) = b(:)
c(1:n) = d(2:n+1)
```

- The `forall` keyword also indicates an array assignment:

```
forall (i=1:n)
    a(i) = b(i)
    c(i) = d(i+1)
end forall
```

You can tell that this is for arrays only, because the loop index has to be part of the left-hand side of every assignment.

This mechanism is prone to misunderstanding and therefore now deprecated. It is not a parallel loop! For that, the following mechanism is preferred.

- The `do concurrent` is a true do-loop. With the `concurrent` keyword the user specifies that the iterations of a loop are independent, and can therefore possibly be done in parallel:

```
do concurrent (i=1:n)
    a(i) = b(i)
    c(i) = d(i+1)
end do
```

#### 34.1.1 Loops without dependencies

Here are some illustrations of simple array copying with the above mechanisms.

```
do i=2,n
    counted(i) = 2*counting(i-1)
end do
```

Original	1	2	3	4	5	6	7	8	9	10
Recursive	0	2	4	6	8	10	12	14	16	18

```
counted(2:n) = 2*counting(1:n-1)
```

Original	1	2	3	4	5	6	7	8	9	10
Sliced	0	2	4	6	8	10	12	14	16	18

```
forall (i=2:n)
    counted(i) = 2*counting(i-1)
end forall
```

Original	1	2	3	4	5	6	7	8	9	10
Forall	0	2	4	6	8	10	12	14	16	18

```
do concurrent (i=2:n)
    counted(i) = 2*counting(i-1)
end do
```

Original	1	2	3	4	5	6	7	8	9	10
Concurrent	0	2	4	6	8	10	12	14	16	18

**Exercise 34.1.** Create arrays A, C of length  $2N$ , and B of length  $N$ . Now implement

$$B_i = (A_{2i} + A_{2i+1})/2, \quad i = 1, \dots, N$$

and

$$C_i = A_{i/2}, \quad i = 1, \dots, 2N$$

using all four mechanisms. Make sure you get the same result every time.

### 34.1.2 Loops with dependencies

For parallel execution of a loop, all iterations have to be independent. This is not the case if the loop has a *recurrence*, and in this case, the ‘do concurrent’ mechanism is not appropriate. Here are the above four constructs, but applied to a loop with a dependence.

```
do i=2,n
    counting(i) = 2*counting(i-1)
end do
```

Original	1	2	3	4	5	6	7	8	9	10
Recursive	1	2	4	8	16	32	64	128	256	512

The slicing version of this:

```
counting(2:n) = 2*counting(1:n-1)
```

Original	1	2	3	4	5	6	7	8	9	10
Sliced	1	2	4	6	8	10	12	14	16	18

acts as if the right-hand side is saved in a temporary array, and subsequently assigned to the left-hand side.

Using ‘forall’ is equivalent to slicing:

```
forall (i=2:n)
    counting(i) = 2*counting(i-1)
end forall
```

Original	1	2	3	4	5	6	7	8	9	10
Forall	1	2	4	6	8	10	12	14	16	18

On the other hand, ‘do concurrent’ does not use temporaries, so it is more like the sequential version:

```
do concurrent (i=2:n)
    counting(i) = 2*counting(i-1)
end do
```

Original	1	2	3	4	5	6	7	8	9	10
Concurrent	1	2	4	8	16	32	64	128	256	512

Note that the result does not have to be equal to the sequential execution: the compiler is free to rearrange the iterations any way it sees fit.



**PART IV**

**EXERCISES AND PROJECTS**



## Chapter 35

### Simple exercises

#### 35.1 Looping exercises

Exercise 35.1. Find all triples of integers  $u, v, w$  under 100 such that  $u^2 + v^2 = w^2$ . Make sure you omit duplicates of solutions you have already found.

Exercise 35.2. The integer sequence

$$u_{n+1} = \begin{cases} u_n/2 & \text{if } u_n \text{ is even} \\ 3u_n + 1 & \text{if } u_n \text{ is odd} \end{cases}$$

leads to the Collatz conjecture: no matter the starting guess  $u_1$ , the sequence  $n \mapsto u_n$  will always terminate.

For  $u_1 < 1000$  find the values that lead to the longest sequence: every time you find a sequence that is longer than the previous maximum, print out the starting number.

Exercise 35.3. Large integers are often printed with a comma (US usage) or a period (European usage) between all triples of digits. Write a program that accepts an integer such as 2542981 from the user, and prints it as 2, 542, 981.

Exercise 35.4. **Root finding.** For many functions  $f$ , finding their zeros, that is, the values  $x$  for which  $f(x) = 0$ , can not be done analytically. You then have to resort to numerical *root finding* schemes. In this exercise you will implement a simple scheme.

Suppose  $x_-, x_+$  are such that

$$x_- < x_+, \quad f(x_-) \cdot f(x_+) < 0,$$

that is, the function values are of opposite sign. Then there is a zero in the interval  $(x_-, x_+)$ . Try to find this zero by looking at the halfway point, and based on the function value there, considering the left or right interval.

- How do you find  $x_-, x_+$ ? This is tricky in general; if you can find them in the interval  $[-10^6, +10^6]$ , halt the program.
- Finding the zero exactly may also be tricky or maybe impossible. Stop your program if  $|x_- - x_+| < 10^{-10}$ .

### 35.1.1 Further practice

The website <http://www.codeforwin.in/2015/06/for-do-while-loop-programming-exercises.html> lists the following exercises:

1. Write a C program to print all natural numbers from 1 to n. - using while loop
2. Write a C program to print all natural numbers in reverse (from n to 1). - using while loop
3. Write a C program to print all alphabets from a to z. - using while loop
4. Write a C program to print all even numbers between 1 to 100. - using while loop
5. Write a C program to print all odd number between 1 to 100.
6. Write a C program to print sum of all even numbers between 1 to n.
7. Write a C program to print sum of all odd numbers between 1 to n.
8. Write a C program to print table of any number.
9. Write a C program to enter any number and calculate sum of all natural numbers between 1 to n.
10. Write a C program to find first and last digit of any number.
11. Write a C program to count number of digits in any number.
12. Write a C program to calculate sum of digits of any number.
13. Write a C program to calculate product of digits of any number.
14. Write a C program to swap first and last digits of any number.
15. Write a C program to find sum of first and last digit of any number.
16. Write a C program to enter any number and print its reverse.
17. Write a C program to enter any number and check whether the number is palindrome or not.
18. Write a C program to find frequency of each digit in a given integer.
19. Write a C program to enter any number and print it in words.
20. Write a C program to print all ASCII character with their values.
21. Write a C program to find power of any number using for loop.
22. Write a C program to enter any number and print all factors of the number.
23. Write a C program to enter any number and calculate its factorial.
24. Write a C program to find HCF (GCD) of two numbers.
25. Write a C program to find LCM of two numbers.
26. Write a C program to check whether a number is Prime number or not.
27. Write a C program to check whether a number is Armstrong number or not.
28. Write a C program to check whether a number is Perfect number or not.
29. Write a C program to check whether a number is Strong number or not.
30. Write a C program to print all Prime numbers between 1 to n.
31. Write a C program to print all Armstrong numbers between 1 to n.
32. Write a C program to print all Perfect numbers between 1 to n.
33. Write a C program to print all Strong numbers between 1 to n.
34. Write a C program to enter any number and print its prime factors.
35. Write a C program to find sum of all prime numbers between 1 to n.
36. Write a C program to print Fibonacci series up to n terms.
37. Write a C program to find one's complement of a binary number.
38. Write a C program to find two's complement of a binary number.
39. Write a C program to convert Binary to Octal number system.
40. Write a C program to convert Binary to Decimal number system.
41. Write a C program to convert Binary to Hexadecimal number system.

42. Write a C program to convert Octal to Binary number system.
43. Write a C program to convert Octal to Decimal number system.
44. Write a C program to convert Octal to Hexadecimal number system.
45. Write a C program to convert Decimal to Binary number system.
46. Write a C program to convert Decimal to Octal number system.
47. Write a C program to convert Decimal to Hexadecimal number system.
48. Write a C program to convert Hexadecimal to Binary number system.
49. Write a C program to convert Hexadecimal to Octal number system.
50. Write a C program to convert Hexadecimal to Decimal number system.
51. Write a C program to print Pascal triangle upto n rows.

## 35.2 Object oriented exercises

Exercise 35.5. Use the `Point` class from exercise ... and make a `Set` class, where each `Set` object contains multiple points.

```
class Point {  
private:  
    float x,y;  
public:  
    Point(float ux,float uy) { x = ux; y = uy; };  
    // other methods omitted  
};  
int main() {  
    Set set;  
  
    set.add( Point(1.0,1.0) );  
    set.add( Point(2.0,2.0) );  
    cout << "Number of points: " << set.cardinality() << endl;  
  
    Point p3(4.0,5.0);  
    set.add( p3 );  
    cout << "Number of points: " << set.cardinality() << endl;  
  
    return 0;  
}
```

This program prints out first 2, then 3, as the number of points in the set.



## **Chapter 36**

### **Not-so simple exercises**

#### **36.1 List access**

Exercise 36.1. Explore the efficiency of using an array versus a linked list.

1. Compare re-allocating the array versus adding elements to the linked list. Start with a simple case: add elements only at the end, and keep a pointer to the final element in the list.
2. Investigate access times: retrieve many (as in: thousands if not more) elements from the array. Do this as follows: allocate an array of indexes, and repeatedly retrieve those list/array elements, for instance adding them together. Does the access time for the array go up if the number of elements gets large?
3. Optimize allocation for the list: create an array of list nodes and use those. Does this make a difference in access times?



# Chapter 37

## Prime numbers

### 37.1 Arithmetic

*Before doing this section, make sure you study section 4.4.*

Exercise 37.1. Read two integers into two variables, and print their sum, product, quotient, modulus.

A less common operator is the modulo operator %.

Exercise 37.2. Read two numbers and print out their modulus. Two ways:

- use the cout function to print the expression, or
- assign the expression to a variable, and print that variable.

### 37.2 Conditionals

*Before doing this section, make sure you study section 5.1.*

Exercise 37.3. Read two numbers and print a message like

3 is a divisor of 9

if the first is an exact divisor of the second, and another message

4 is not a divisor of 9

if it is not.

### 37.3 Looping

*Before doing this section, make sure you study section 6.1.*

Exercise 37.4. Read an integer and determine whether it is prime by testing for the smaller numbers whether they are a divisor of that number.

Print a final message

Your number is prime

or

Your number is not prime: it is divisible by ....

where you report just one found factor.

**Exercise 37.5.** Rewrite the previous exercise with a boolean variable to represent the primeness of the input number.

## 37.4 Functions

*Before doing this section, make sure you study section ??.*

Above you wrote several lines of code to test whether a number was prime.

**Exercise 37.6.** Write a function that takes an integer input, and return a boolean corresponding to whether the input was prime.

```
bool isprime;  
isprime = prime_test_function(13);
```

Read the number in, and print the value of the boolean.

## 37.5 While loops

*Before doing this section, make sure you study section 6.2.*

**Exercise 37.7.** Take the prime number testing program, and modify it to read in how many prime numbers you want to print. Print that many successive primes. Keep a variable `number_of_primes_found` that is increased whenever a new prime is found.

## 37.6 Global variables: optional

*Before doing this section, make sure you study section ??.*

**Exercise 37.8.** Use global variables to rewrite exercise ???. Your main program should exactly be this:

```
int main() {  
    int nprimes;  
    cout << "How many primes do you want? " << endl;  
    cin >> nprimes;  
  
    while (numberfound<nprimes) {  
        int number = nextprime();  
        cout << "Number " << number << " is prime" << endl;  
    }  
  
    return 0;  
}
```

The trick here is to write the function `nextprime` uses the remembered global information, calculates the next prime, and returns it.

## 37.7 Structures

*Before doing this section, make sure you study section 9.1, 14.1.*

A struct functions to bundle together a number of data item. We only discuss this as a preliminary to classes.

**Exercise 37.9.** Rewrite the exercise that found a predetermined number of primes, putting the `number_of_primes_found` and `last_number_tested` variables in a structure. Your main program should now look like:

```
struct primesequence sequence;
while (sequence.number_of_primes_found < nprimes) {
    int number = nextprime(sequence);
    cout << "Number " << number << " is prime" << endl;
}
```

## 37.8 Classes and objects

*Before doing this section, make sure you study section 10.1, 30.1.*

In exercise ?? you made a structure that contains the data for a primesequence, and you have separate functions that operate on that structure or on its members.

**Exercise 37.10.** Write a class `primesequence` that contains the members of the structure, and the functions `nextprime`, `isprime`. The function `nextprime` does not need the structure as argument, because the structure members are in the class, and therefore global to that function.

Your main program should look as follows:

```
primesequence sequence;
while (sequence.numberfound < nprimes) {
    int number = sequence.nextprime();
    cout << "Number " << number << " is prime" << endl;
}
```

In the previous exercise you defined the `primesequence` class, and you made one object of that class:

```
primesequence sequence;
```

But you can make multiple sequences, that all have their own internal data and are therefore independent of each other.

**Exercise 37.11.** The *Goldbach conjecture* says that every even number, from 4 on, is the sum of two primes  $p+q$ . Write a program to test this for the even numbers up to 20 million. Make an outer loop over the even numbers. In each iteration, make a `primesequence` object to generate  $p$  values. Then, for each  $p$ , make a second `primesequence` object to generate  $q$  values, and test with these.

For each even number, print out how it is the sum of two primes. If multiple possibilities exist, only print the first one you find.

**Exercise 37.12.** The *Goldbach conjecture* says that every even number  $2n$  (starting at 4), is the sum of two primes  $p + q$ :

$$2n = p + q.$$

Equivalently, every number  $n$  is equidistant from two primes. In particular this holds for each prime number:

$$\forall_{p \text{ prime}} \exists_{q \text{ prime}} : r \equiv p + (p - q) \text{ is prime.}$$

Write a program that tests this. You need two prime number generators, one for the  $p$ -sequence and one for the  $q$ -sequence. For each  $p$  value, when the program finds the  $q$  value, print the  $q, p, r$  triple and move on to the next  $p$ .

Allocate an array where you record all the  $p - q$  distances that you found. Print some elementary statistics, for instance: what is the average, do the distances increase or decrease with  $p$ ?

## 37.9 Arrays

Another algorithm for finding prime numbers is the *Eratosthenes sieve*. It goes like this.

1. You take a range of integers, starting at 2.
2. Now look at the first number. That's a prime number.
3. Scratch out all of its multiples.
4. Find the next number that's not scratched out; since that's not a multiple of a previous number, it must be a prime number. Report it, and go back to the previous step.

The new mechanism you need for this is the data structure for storing all the integers.

```
int N = 1000;
vector<int> integers(N);
```

**Exercise 37.13.** Read in an integer that denotes the largest number you want to test. Make an array of integers that long. Set the elements to the successive integers.

## Chapter 38

### Infectuous disease simulation

This section contains a sequence of exercises that builds up to a somewhat realistic simulation of the spread of infectious diseases. The main skill exercised here is the use of arrays; section 11. You are greatly encouraged to use *separate compilation* and *header files*; section 18.1.1.

#### 38.1 Basic design

*Before doing this section, make sure you study section 11.1.*

It is possible to model disease propagation statistically, but here we will build an explicit simulation: we will maintain an array of all the people in the population, and track for each of them their status. That status is modeled by a single integer, and it can be:

- inoculated or recovered, which we model by  $-1$ ,
- healthy but not inoculated, which we model by  $0$ ,
- and sick, with  $n$  days to go before recovery, which we model by  $n$ .

Any sick person is infectious, and we will look at different ways of modeling how the disease spreads from an infectious person. The program will then track the population from day to day, running indefinitely until none of the population is sick. (Since there is no re-infection, the run will always end.)

Exercise 38.1. Start setting up the basic program design. Use one source file for the main program and one for utility functions. Make sure to use a header file (see section 18.1.1).

- Write a function that, given a person's status on one day (meaning: an integer) computes their status on the next day.
- Write a function that gives a random disease duration, make this a number between 5 and 10.
- Create an array representing the population, where you read the size from terminal input. Mark just one person as being sick, everyone else is healthy but not inoculated.
- Now set up the main loop: on each day (meaning: in each step) you
  - give every person a new status
  - count how many people are sick. If this number drops to zero, end the simulation.

For the state of the population you can use a static array. Write a function that displays the state of the population on a given day.

This past exercise was too simplistic: the original patient zero was the only one who ever got sick. Now we incorporate contagion.

**Exercise 38.2.** Read in a number  $0 \leq p \leq 1$  representing the probability of disease transmission upon contact. Incorporate this into the program: in each step the direct neighbours of an infected person get sick themselves.

You need to think about program design here: can you do the updates directly in the population array?

Run a number of simulations with population sizes and contagion probabilities. Are there cases where people escape getting sick?

**Exercise 38.3.** Incorporate inoculation: read another number representing the percentage of people that has been vaccinated. Choose those members of the population randomly. Describe the effect of vaccinated people on the spread of the disease. Why is this model unrealistic.

To make the simulation more realistic, we let every sick person come into contact with a fixed number of random people every day. Set this number to 6 or so. Use another probability for the chance that the sick person actually transmits the disease.

**Exercise 38.4.** Code the random interactions. Now run a number of simulations varying

- The percentage of people inoculated, and
- the chance the disease is transmitted on contact.

Record how long the disease runs through the population. With a fixed degree of contagiousness, how is this number of function of the percentage that is vaccinated?

Investigate the matter of ‘herd immunity’: if enough people are vaccinated, then some people who are not will still never get sick. Let’s say you want to have this probability over 95 percent. Investigate the percentage of inoculation that is needed for this as a function of the contagiousness of the disease.

## Chapter 39

### Geometry

In this set of exercises you will write a small ‘geometry’ package: code that manipulates points, lines, shapes. These exercises mostly use the material of section 10.

#### 39.1 Point class

*Before doing this section, make sure you study section 10.1.*

A class can contain elementary data. In this section you will make a `Point` class that models cartesian coordinates and functions defined on coordinates.

Exercise 39.1. Make class `Point` with a constructor

```
Point( float xcoordinate, float ycoordinate );
```

Write a method `distance_to_origin` that returns a `float`.

Write a method `printout` that uses `cout` to display the point.

Write a function `distance` so that if `p, q` are `Point` objects,

```
p.distance(q)
```

computes the distance.

Exercise 39.2. Make a default constructor for the point class:

```
Point() { /* default code */ }
```

which you can use as:

```
Point p;
```

Exercise 39.3. Advanced. Can you make a `Point` class that can accomodate any number of space dimensions? Hint: use a `vector`; section 11.4. Can you make a constructor where you do not specify the space dimension explicitly?

#### 39.2 Using one class in another

*Before doing this section, make sure you study section 10.2.*

Exercise 39.4. Make a class `LinearFunction` with a constructors:

```
LinearFunction( Point input_p1,Point input_p2 );
```

and a function

```
float evaluate_at( float x );
```

which you can use as:

```
LinearFunction line(p1,p2);
cout << "Value at 4.0: " << line.evaluate_at(4.0) << endl;
```

**Exercise 39.5.** Make a class `LinearFunction` with two constructors:

```
LinearFunction( Point input_p2 );
LinearFunction( Point input_p1,Point input_p2 );
```

where the first stands for a line through the origin.

Implement again the `evaluate` function so that

```
LinearFunction line(p1,p2);
cout << "Value at 4.0: " << line.evaluate_at(4.0) << endl;
```

Suppose you want to write a `Rectangle` class, which could have methods such as `float Rectangle::area()` or `bool Rectangle::contains(Point)`. Since rectangle has four corners, you could store four `Point` objects in each `Rectangle` object. However, there is redundancy there: you only need three points to infer the fourth. Let's consider the case of a rectangle with sides that are horizontal and vertical; then you need only two points.

Intended API:

```
float Rectangle::area();
```

It would be convenient to store width and height; for

```
bool Rectangle::contains(Point);
```

it would be convenient to store bottomleft/topright points.

**Exercise 39.6.** Make a class `Rectangle` (sides parallel to axes) with two constructors:

```
Rectangle(Point bl,Point tr);
Rectangle(Point bl,float w,float h);
```

and functions

```
float area(); float width(); float height();
```

Let the `Rectangle` object store two `Point` objects.

Then rewrite your exercise so that the `Rectangle` stores only one point (say, lower left), plus the width and height.

The previous exercise illustrates an important point: for well designed classes you can change the implementation (for instance motivated by efficiency) while the program that uses the class does not change.

### 39.3 Is-a relationship

*Before doing this section, make sure you study section 10.3.*

Exercise 39.7. Take your code where a `Rectangle` was defined from one point, width, and height.

Make a class `Square` that inherits from `Rectangle`. It should have the function `area` defined, inherited from `Rectangle`.

First ask yourself: what should the constructor of a `Square` look like?

Exercise 39.8. Revisit the `LinearFunction` class. Add methods `slope` and `intercept`.

Now generalize `LinearFunction` to `StraightLine` class. These two are almost the same except for vertical lines. The `slope` and `intercept` do not apply to vertical lines, so design `StraightLine` so that it stores the defining points internally. Let `LinearFunction` inherit.



# Chapter 40

## PageRank

### 40.1 Basic ideas

Assuming you have learned about arrays 11, in particular the use of `std::vector`.

The web can be represented as a matrix  $W$  of size  $N$ , the number of web pages, where  $w_{ij} = 1$  if page  $i$  has a link to page  $j$  and zero otherwise. However, this representation is only conceptual; if you actually stored this matrix it would be gigantic and largely full of zeros. Therefore we use a *sparse matrix*: we store only the pairs  $(i, j)$  for which  $w_{ij} \neq 0$ . (In this case we can get away with storing only the indices; in a general sparse matrix you also need to store the actual  $w_{ij}$  value.)

Exercise 40.1. Store the sparse matrix representing the web as a

```
vector< vector<bool> >
```

structure.

1. At first, assume that the number of web pages is given and reserve the outer vector. Read in values for nonzero indices and add those to the matrix structure.
2. Then, assume that the number of pages is not pre-determined. Read in indices; now you need to create rows as they are needed. Suppose the requested indices are

```
5, 1  
3, 5  
1, 3
```

Since your structure has only three rows, you also need to remember their row numbers.

Now we want to model the behaviour of a person clicking on links.



**PART V**

**ADVANCED TOPICS**



# Chapter 41

## Tiniest of introductions to algorithms and data structures

### 41.1 Data structures

This really goes beyond this book.

- Arrays
- Lists
- Trees
- Graphs / DAGs
- Hashes

#### 41.1.1 Linked lists

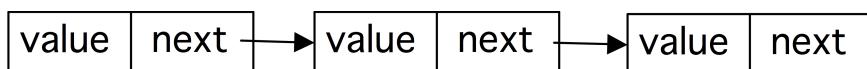
*Before doing this section, make sure you study section 17.*

Arrays are not flexible: you can not insert an element in the middle. Instead:

- Allocate larger array,
- copy data over (with insertion),
- delete old array storage

This is very expensive. (It's what happens in a C++ `vector`; section 11.4.2.)

If you need to do lots of insertions, make a *linked list*.



```
class Node {  
private:  
    int data{0}, count{0};  
    Node *next=nullptr;  
public:  
    Node() {}  
    Node(int value) { data = value; count++; };  
    bool hasnext() {  
        return next!=nullptr; };
```

```
Node *insert(int value) {
    if (value==this->data) {
        // we have already seen this value: just count
        count++;
        return this;
    } else if (value>this->data) {
        // value belong in the tail
        if (!hasnext())
            next = new Node(value);
        else
            next = next->insert(value);
        return this;
    } else {
        // insert at the head of the list
        Node *newhead = new Node(value);
        newhead->next = this;
        return newhead;
    }
};
```

### 41.1.2 Trees

*Before doing this section, make sure you study section 17.*

A tree can be defined recursively:

- A tree is empty, or
- a tree is a node with some number of children trees.

Let's design a tree that stores and counts integers: each node has a label, namely an integer, and a count value that records how often we have seen that integer.

Our basic data structure is the node, and we define it recursively to have up to two children. This is a problem: you can not write

```
class Node {
private:
    Node left,right;
}
```

because that would recursively need infinite memory. So instead we use pointers.

```
class Node {
private:
    int key{0},count{0};
    std::shared_ptr<Node> left,right;
    bool hasleft{false},hasright{false};
public:
    Node() { }
```

```

Node(int i,int init=1) { key = i; count = 1; }
void addleft( int value) {
    left = std::shared_ptr<Node>( new Node(value) );
    hasleft = true;
}
void addright( int value ) {
    right = std::shared_ptr<Node>( new Node(value) );
    hasright = true;
}
};

```

and we record that we have seen the integer zero zero times.

Algorithms on a tree are typically recursive. For instance, the total number of nodes is computed from the root. At any given node, the number of nodes of that attached subtree is one plus the number of nodes of the left and right subtrees.

```

int number_of_nodes() {
    int count = 1;
    if (hasleft)
        count += left->number_of_nodes();
    if (hasright)
        count += right->number_of_nodes();
    return count;
}

```

Likewise, the depth of a tree is computed as a recursive max over the left and right subtrees:

```

int depth() {
    int d = 1, dl=0,dr=0;
    if (hasleft)
        dl = left->depth();
    if (hasright)
        dr = right->depth();
    d = max(d+dl,d+dr);
    return d;
}

```

Now we need to consider how actually to insert nodes. We write a function that inserts an item at a node. If the key of that node is the item, we increase the value of the counter. Otherwise we determine whether to add the item in the left or right subtree. If no such subtree exists, we create it; otherwise we descend in the appropriate subtree, and do a recursive insert call.

```

void insert(int value) {
    if (key==value)
        count++;
    else if (value<key) {
        if (hasleft)

```

```
    left->insert(value);
        else
    addleft(value);
} else if (value>key) {
    if (hasright)
right->insert(value);
    else
addrightright(value);
} else throw(1); // should not happen
};
```

## 41.2 Algorithms

This *really really* goes beyond this book.

- Simple ones: numerical
- Connected to a data structure: search

### 41.2.1 Sorting

Unlike the tree algorithms above, which used a non-obvious data structure, sorting algorithms are a good example of the combination of very simple data structures (mostly just an array), and sophisticated analysis of the algorithm behaviour. We very briefly discuss two algorithms.

#### 41.2.1.1 Bubble sort

An array  $a$  of length  $n$  is sorted if

$$\forall_{i < n-1}: a_i \leq a_{i+1}.$$

A simple sorting algorithm suggests itself immediately: if  $i$  is such that  $a_i > a_{i+1}$ , then reverse the  $i$  and  $i + 1$  locations in the array.

```
void swap( std::vector<int> &array, int i ) {
    int t = array[i];
    array[i] = array[i+1];
    array[i+1] = t;
}
```

(Why is the array argument passed by reference?)

If you go through the array once, swapping elements, the result is not sorted, but at least the largest element is at the end. You can now do another pass, putting the next-largest element in place, and so on.

This algorithm is known as *bubble sort*. It is generally not considered a good algorithm, because it has a time complexity (section 43.1.1) of  $n^2/2$  swap operations. Sorting can be shown to need  $O(n \log n)$  operations, and bubble sort is far above this limit.

### 41.2.1.2 Quicksort

A popular algorithm that can attain the optimal complexity (but need not; see below) is *quicksort*:

- Find an element, called the pivot, that is approximately equal to the median value.
- Rearrange the array elements to give three sets, consecutively stored: all elements less than, equal, and greater than the pivot respectively.
- Apply the quicksort algorithm to the first and third subarrays.

This algorithm is best programmed recursively, and you can even make a case for its parallel execution: every time you find a pivot you can double the number of active processors.

**Exercise 41.1.** Suppose that, by bad luck, your pivot turns out to be the smallest array element every time. What is the time complexity of the resulting algorithm?

## 41.3 Programming techniques

### 41.3.1 Memoization

In section 7.2 you saw some examples of recursion. The factorial example could be written in a loop, and there are both arguments for and against doing so.

The Fibonacci example is more subtle: it can not immediately be converted to an iterative formulation, but there is a clear need for eliminating some waste that comes with the simple recursive formulation. The technique we can use for this is known as *memoization*: store intermediate results to prevent them from being recomputed.

Here is an outline.

```

int fibonacci(int n) {
    std::vector<int> fibo_values(n);
    for (int i=0; i<n; i++)
        fibo_values[i] = 0;
    fibonacci_memoized(fibo_values, n-1);
    return fibo_values[n-1];
}
int fibonacci_memoized( std::vector<int> &values, int top ) {
    int minus1 = top-1, minus2 = top-2;
    if (top<2)
        return 1;
    if (values[minus1]==0)
        values[minus1] = fibonacci_memoized(values, minus1);
    if (values[minus2]==0)
        values[minus2] = fibonacci_memoized(values, minus2);
    values[top] = values[minus1]+values[minus2];
    //cout << "set f(" << top << ") to " << values[top] << endl;
    return values[top];
}
//codesnipet end

```

```
int main() {
    int fibo_n;
    cout << "What number? ";
    cin >> fibo_n;
    cout << "Fibo(" << fibo_n << ") = " << fibonacci(fibo_n) << endl;

    return 0;
}
```

## Chapter 42

### Programming strategies

#### 42.1 Programming: top-down versus bottom up

The exercises in chapter 37 were in order of increasing complexity. You can imagine writing a program that way, which is formally known as *bottom-up* programming.

However, to write a sophisticated program this way you really need to have an overall conception of the structure of the whole program.

Maybe it makes more sense to go about it the other way: start with the highest level description and gradually refine it to the lowest level building blocks. This is known as *top-down* programming.

<https://www.cs.fsu.edu/~myers/c++/notes/stepwise.html>

Example:

Run a simulation

becomes

Run a simulation:

    Set up data and parameters

    Until convergence:

        Do a time step

becomes

Run a simulation:

    Set up data and parameters:

        Allocate data structures

        Set all values

    Until convergence:

        Do a time step:

            Calculate Jacobian

            Compute time step

            Update

You could do these refinement steps on paper and wind up with the finished program, but every step that is refined could also be a subprogram.

We already did some top-down programming, when the prime number exercises asked you to write functions and classes to implement a given program structure; see for instance exercise ??.

A problem with top-down programming is that you can not start testing until you have made your way down to the basic bits of code. With bottom-up it's easier to start testing. Which brings us to...

### 42.1.1 Worked out example

Take a look at exercise 35.2. We will solve this in steps.

1. State the problem:

```
// find the longest sequence
```

2. Refine by introducing a loop

```
// find the longest sequence:
```

```
// Try all starting points
```

```
// If it gives a longer sequence report
```

3. Introduce the actual loop:

```
// Try all starting points
```

```
for (int starting=2; starting<1000; starting++)
```

```
// If it gives a longer sequence report
```

4. Record the length:

```
// Try all starting points
```

```
int maximum_length=-1;
```

```
for (int starting=2; starting<1000; starting++) {
```

```
// If the sequence gives a longer sequence report:
```

```
int length=0;
```

```
// compute the sequence
```

```
if (length>maximum_length) {
```

```
// Report this sequence as the longest
```

```
}
```

```
}
```

5. // Try all starting points

```
int maximum_length=-1;
```

```
for (int starting=2; starting<1000; starting++) {
```

```
// If the sequence gives a longer sequence report:
```

```
int length=0;
```

```
// compute the sequence
```

```
int current=starting;
```

```
while (current!=1) {
```

```
// update current value
```

```
length++;
```

```
}
```

```
if (length>maximum_length) {
```

```
// Report this sequence as the longest
```

```
}
```

```
}
```

## 42.2 Coding style

After you write your code there is the issue of *code maintainance*: you may in the future have to update your code or fix something. You may even have to fix someone else's code or someone will have to work on your code. So it's a good idea to code cleanly.

**Naming** Use meaningful variable names: `record_number` instead `rn` or `n`. This is sometimes called 'self-documenting code'.

**Comments** Insert comments to explain non-trivial parts of code.

**Reuse** Do not write the same bit of code twice: use macros, functions, classes.

## 42.3 Documentation

Take a look at Doxygen.

## 42.4 Testing

If you write your program modularly, it is easy (or at least: easier) to test the components without having to wait for an all-or-nothing test of the whole program. In an extreme form of this you would write your code by *test-driven development*: for each part of the program you would first write the test that it would satisfy.

In a more moderate approach you would use *unit testing*: you write a test for each program bit, from the lowest to the highest level.

And always remember the old truism that 'by testing you can only prove the presence of errors, never the absence.'



## **Chapter 43**

### **Complexity**

#### **43.1 Order of complexity**

##### **43.1.1 Time complexity**

Exercise 43.1. For each number  $n$  from 1 to 100, print the sum of all numbers 1 through  $n$ .

There are several possible solutions to this exercise. Let's assume you don't know the formula for the sum of the numbers  $1 \dots n$ . You can have a solution that keeps a running sum, and a solution with an inner loop.

Exercise 43.2. How many operations, as a function of  $n$ , are performed in these two solutions?

##### **43.1.2 Space complexity**

Exercise 43.3. Read numbers that the user inputs; when the user inputs zero or negative, stop reading. Add up all the positive numbers and print their average.

This exercise can be solved by storing the numbers in a `std::vector`, but one can also keep a running sum and count.

Exercise 43.4. How much space do the two solutions require?



## Chapter 44

### C for C++ programmers

Youll have no exceptions, so unwinding the stack on error is manual. Cue lots of if (retval == NULL) stuff - most projects have a common idiom. Mind you get everything, since missing one error case will probably cause either a crash or a security nightmare, or possibly both.

Youll have no destructors, so cleanup is manual. This is most fun with early-return functions, but it can keep you entertained for all cases. File handles, memory, and other resources (thread locks, anyone) are all waiting patiently and silently for you to forget them.

Initialization has to be explicitly called. No constructors either.

Want inheritance? Sure. Write your own vtable (often done with function pointers in a struct). Instead of templates, youll need to abandon type safety and cast back and forth to (void\*). Dont explicitly cast to (void \*), because the compiler never warns about explicit or implicit casts to and from (void \*).

Youll also need to make sure youre using the right library calls - snprintf versus sprintf etc. Hopefully an existing project will be using the right ones.

On the plus side, youre moving to Linux, and a lot of the tooling available is - while very command-line oriented - very good indeed.

For an IDE, Id recommend CLion from JetBrains, but Im told that with sufficient patience, Atom can be encouraged into doing useful stuff.

Youll find that while the command-line of GDB, the debugger, isnt very easy to learn to begin with, its very powerful, allowing you to do conditional breakpoints with comparative ease.

Valgrind is amazing. Voodoo. Itll find uninitialized memory, allocation errors, overflows, and leaks - all common and hard to debug issues in C.

The CLang static analyzer is pretty impressive, too.

(copied from <https://www.quora.com/How-should-a-C--programmer-learn-Linux-C/answer/Dave-Cridland>)



## **PART VI**

### **INDEX AND SUCH**



## Chapter 45

### Index

accessor, 54  
allocate, 140  
Apple, 17  
argument  
    default, 47  
array  
    automatic, 139  
    static, 139  
assignment, 29  
  
bad\_alloc, 109  
bad\_exception, 109  
bit\_size, 123  
bottom-up, 183  
break, 39  
bubble sort, 69, 180  
  
C preprocessor, see preprocessor  
C++, 116  
C++11, 64, 91  
    range-based iterator, 112  
c(sizeof, 123  
call, 128  
calling environment, 79  
case sensitive, 28  
cast, 33, 113, 113  
cin, 31  
class  
    abstract, 60  
    base, 59  
    derived, 58  
class, 53, 137  
close, 150  
code  
    maintainance, 185  
  
code reuse, 42  
column-major, 149  
compilation  
    separate, 167  
compiler, 18  
    and preprocessor, 101  
compiling, 18  
conditional, 35  
constructor, 53  
container, 111  
contains  
    for class functions, 133  
    in modules, 129  
contains, 127  
continuation character, 122  
continue, 39  
cout, 30  
Cshift, 141  
  
datatype, 28  
deallocate, 140  
dereference, 86  
destructor  
    at end of scope, 49  
do concurrent, 151  
do loop  
    implicit, 149  
DotProduct, 141  
  
emacs, 17, 17, 121  
end, 121  
Eratosthenes sieve, 166  
exception  
    catch, 108  
    throwing, 107

executable, 18  
exit, 125  
expression, 29  
F90, 19  
false, 30, 32  
floating point, 30  
forall, 151  
Fortran  
    comments, 122  
    source format  
        fixed, 122  
        free, 122  
forward declaration, 93  
function, 41, 128  
    arguments, 43  
    body, 43  
    call, 42  
    defines scope, 44  
    definition, 42  
    parameters, 43  
    result type, 43  
functional programming, 44  
  
g++, 18  
getline, 31  
gfortran, 19  
GNU, 18  
Goldbach conjecture, 165, 166  
  
has-a relation, 58  
header, 167  
header file, 94, 101  
    and global variables, 95  
    treatment by preprocessor, 95  
hexadecimal, 85  
homebrew, 17  
  
icpc, 18  
ifort, 19  
inheritance, 58  
initialization  
    variable, 28  
inline, 129  
is-a relation, 58  
is\_eof, 78  
is\_open, 78  
  
iso\_c\_binding, 123  
iterator, 111  
keywords, 27  
label, 150  
linked list, 145  
Linux, 17  
list  
    linked, 177  
loop, 37  
    body, 37  
    counter, 37  
    for, 37  
    header, 37  
    variable, 37  
    while, 37  
lvalue, 114  
  
macports, 17  
makefile, 94  
malloc, 64, 90  
MatMul, 141  
matrix  
    sparse, 173  
MaxLoc, 141  
MaxVal, 141  
member (of struct), 51  
memoization, 181  
memory leaking, 90  
method  
    abstract, 60  
    overriding, 59  
method, 54  
methods (of an object), 53  
Microsoft  
    Windows, 17  
    Word, 12  
MinLoc, 141  
MinVal, 141  
move semantics, 116  
  
namespace, 97  
new, 63, 89  
Newton's method, 44  
  
object

destructor, 83  
    open, 150  
operators  
    shortcut, 32  
overloading, 57  
  
package manager, 17  
parameter, see function, parameter  
    passing  
        by reference, 79  
        by value, 79  
pass by reference, 45  
pass by value, 44  
pointer  
    arithmetic, 88  
    void, 113  
pointer, 143  
preprocessor  
    and header files, 95  
    conditionals, 102–103  
    macro  
        parametrized, 102  
    macros, 101–102  
print, 149  
private, 54  
procedures  
    internal, 128  
    module, 128  
Product, 141  
program  
    statements, 18  
protected, 59  
prototype, 93  
public, 54  
putty, 17  
python, 15  
quicksort, 181  
  
recurrence, 152  
reference  
    const  
        to class member, 80  
    to class member, 80  
return, 43  
return, 128  
root finding, 157  
  
    rvalue, 114  
    reference, 116  
scope, 43  
    dynamic, 49  
    in conditional branches, 36  
    lexical, 49  
    of function body, 44  
    out of, 83  
size, 140  
size\_t, 113  
sizeof, 90  
smartphone, 12  
source code, 18  
stack, 62, 140  
    overflow, 140  
Standard Template Library, 111  
statement functions, 129  
static\_cast, 113  
storage\_size, 123  
string, 71  
    concatenation, 71  
    size, 71  
struct, 51  
subprogram, see function  
Sum, 141  
swap, 116  
switch, 36  
  
target, 143  
templates  
    and separate compilation, 96  
test-driven development, 185  
testing, 185  
top-down, 183  
Transpose, 141  
true, 30, 32  
type  
    derived, 131  
type, 131  
  
unit testing, 185  
Unix, 17  
use, 129  
  
values  
    boolean, 30

variable, **28**  
    assignment, **28**  
    declaration, **28, 28**  
    global  
        in header file, **95**  
    initialization, **30**  
    numerical, **29**  
    static, **49**  
vector, **97**  
vector, **177**

vi, **17**  
Virtualbox, **17**  
Visual Studio, **17**  
VMware, **17**  
void, **42, 43**  
while, **39**  
Xcode, **17**  
XQuartz, **17**